



Inspiring Excellence

# Application Layer: FTP, P2P and CDN

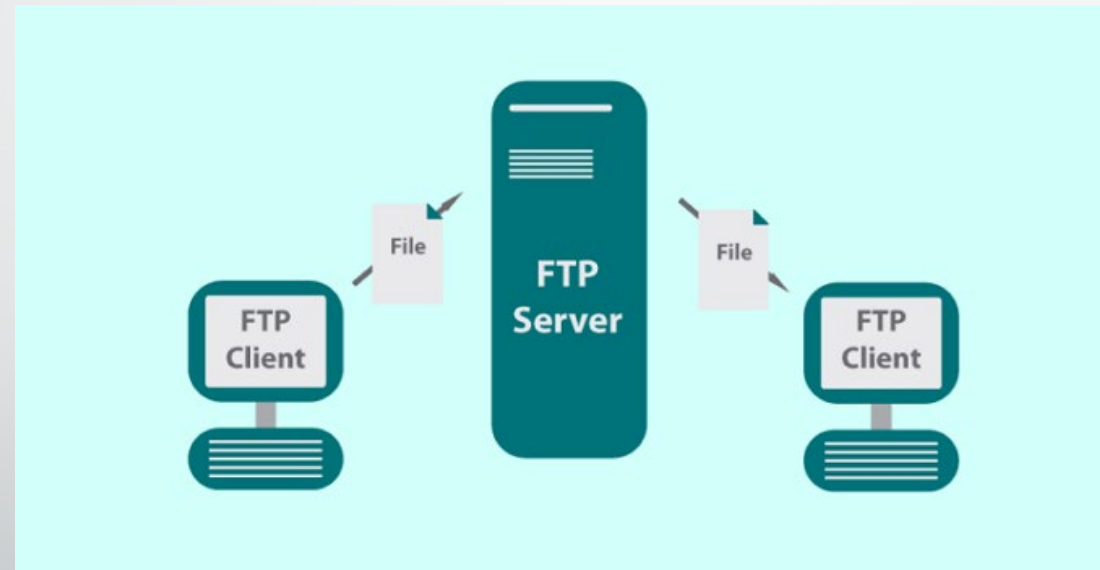
Lecture 5 | CSE421 – Computer Networks

Department of Computer Science and Engineering  
School of Data & Science

# Objectives

- File Distribution
- FTP
- Client-Server Architecture
- P2P Architecture
- CDN

# FTP



# FTP

- FTP stands for **File transfer protocol**.
- FTP is a standard internet protocol **provided by TCP/IP** used for transmitting the **files** from one host to another.
- It is **mainly used for transferring the web page files** from their creator to the computer that acts as a server for other computers on the internet.
- It is also used for **downloading the files to computer** from other servers.
- Objectives:
  - It provides the **sharing** of files.
  - It is used to **encourage** the use of remote computers.
  - It transfers the data more **reliably and efficiently**

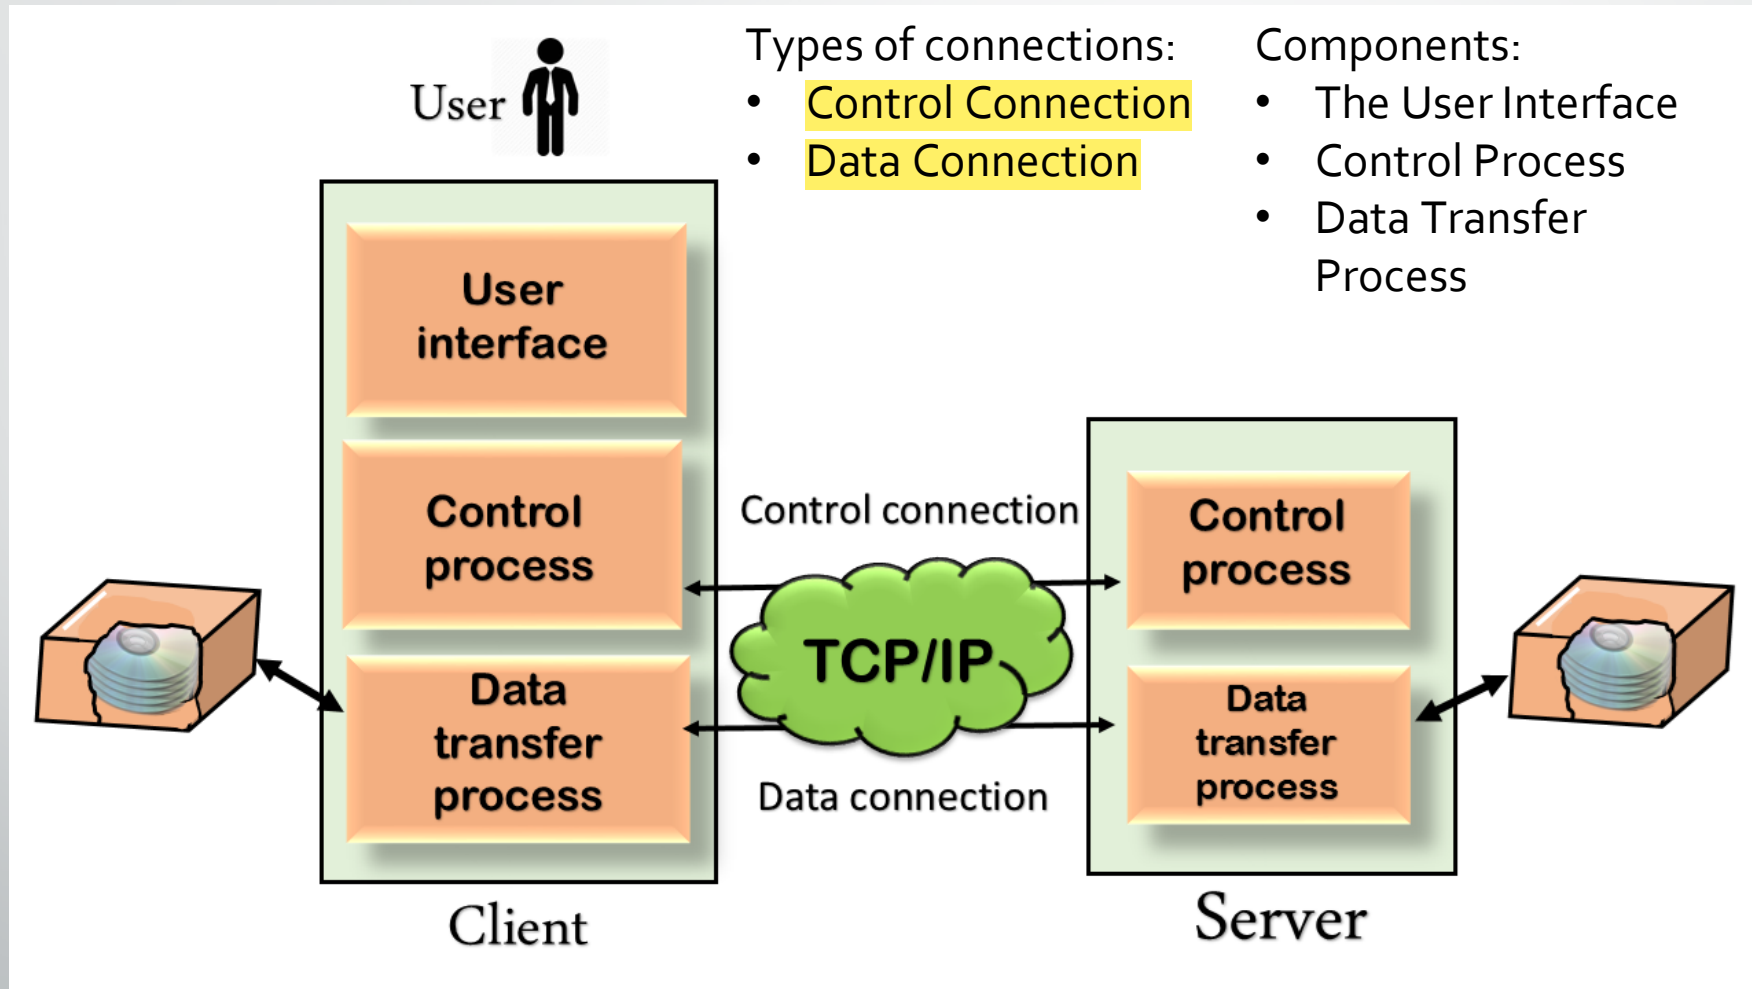
# Why FTP?

Two systems may have different:

- file conventions.
- ways to represent text and data.
- directory structures.

FTP protocol overcomes these problems by establishing two connections between hosts. One connection is used for data transfer, and another connection is used for the control connection.

# Mechanism of FTP



# Types of FTP Connections

- **Control Connection:** The control connection uses very simple rules for communication. Through control connection, we can transfer a line of command or line of response at a time. The control connection is made between the control processes. The control connection **remains connected during the entire interactive FTP session.**
- **Data Connection:** The Data Connection uses very **complex rules** as data types may vary. The data connection is made between data transfer processes. The data connection **opens when a command comes** for transferring the files and **closes when the file is transferred.**

# FTP Clients

- FTP client is a program that **implements a file transfer protocol** which allows you to transfer files between two hosts on the internet.
- It allows a user to **connect to a remote host** and **upload or download** the files.
- It has a **set of commands** that we can use to connect to a host, **transfer** the files between you and your host and **close** the connection.



# FTP Ups and Downs


## Advantages

- **Speed:** The FTP is one of the **fastest way** to transfer the files from one computer to another computer.
- **Efficient:** It is more efficient as we do not need to complete all the operations to get the entire file.
- **Security:** To access the FTP server, we need to login with the username and password.
- **Back & forth movement:** Suppose you are a manager of the company, you send some information to all the employees, and they all send information back on the **same server**.

## Disadvantages

- The standard requirement of the industry is that all the **FTP transmissions should be encrypted**. However, not all the FTP providers are equal and **not all the providers offer encryption**.
- FTP has the **size limit of the file is 2GB** that can be sent. It also **doesn't allow** you to run **simultaneous transfers to multiple receivers**.
- **Passwords and file contents are sent in clear text** that allows **unwanted eavesdropping**.
- It is **not compatible** with every system.

# P2P Applications




# Pure P2P Architecture

- No always-on server
- Arbitrary end systems directly communicate
- Peers are intermittently connected and **change IP addresses**
- Files are **shared in chunks** rather than a whole single file
- A successful file transfer is possible if all clients collectively have all the chunks of a file

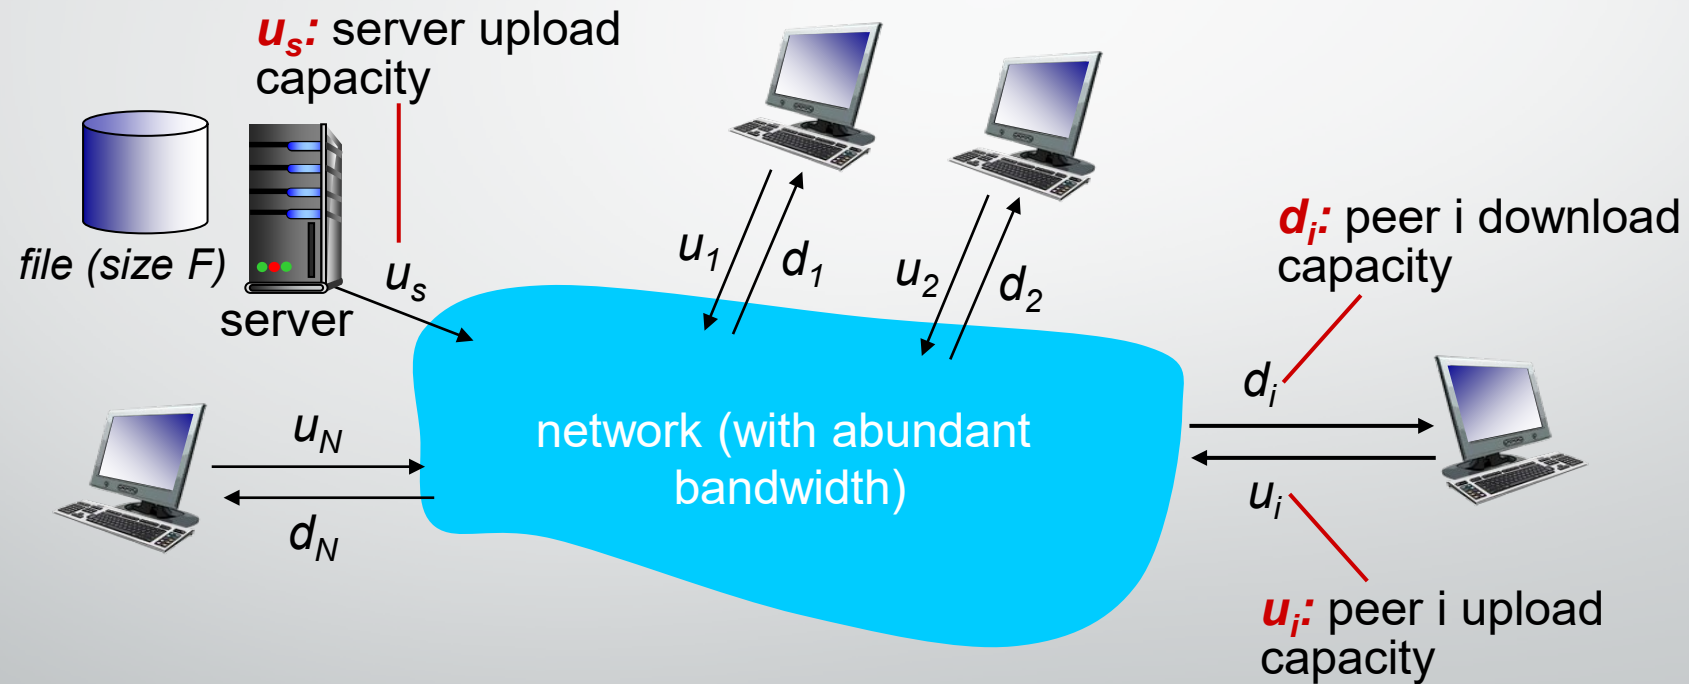
## Examples:

- File distribution (BitTorrent)
- Streaming (KanKan)
- VoIP (Skype)



# File Distribution

- How much time is required to distribute a file (of size  $F$ ) from one server to 'N' number of peers?
  - Peer (client) upload/download capacity is limited resource



# Terms of equations

- Given:
  - File size =  $f$
  - Number of clients =  $n$
  - Server upload speed =  $u_s$
  - Download speed of peer 'i' =  $d_i$
  - Upload speed of peer 'i' =  $u_i$
  - Time to distribute files using client-server approach =  $T_{c-s}$
  - Time to distribute files using peer-to-peer approach =  $T_{P2P}$
  - The peer with the slowest download speed =  $d_{\min}$
  - Summation of upload speed of all peers =  $\sum u_i$

# File Distribution: Client-Server

Server:

- Time to send one copy of file from the server =  $f/u_s$
- Hence, time to **send** this one file to n number of clients =  $n * f/u_s = nf/u_s$

Client:

- **Downloading** time of the slowest client =  $f/d_{\min}$

Time to distribute file 'f' to 'n' clients using client-server approach

- $T_{c-s} \geq \max \{nf/u_s, f/d_{\min}\}$ 
  - So, why the max value of the above two?
    - It's because, if the **server** needs 10 minutes to upload a file, a client can **never** download it before 10 minutes.
    - If the **slowest client** needs 15 minutes to download the file, even though server needs 10 minutes to upload, there's no point. The transfer **won't finish** before 15 minutes.
  - Lastly,  $T_{c-s}$  will take a equal or greater value because, this is the minimum time possible (does not consider any delays). In real world, speed is always not at its maximum, speed varies.

# File Distribution: P2P

**Server transmission:** must upload at least one copy


- Time to upload file  $f$  from the server =  $f/u_s$
- Downloading time of the slowest client =  $f/d_{\min}$
- Total downloaded file size by  $n$  clients =  $n * f = nf$
- The more clients participate in the file sharing.. the more upload speed
- Total upload speed of  $n$  clients =  $u_1 + u_2 + u_3 + \dots + u_n = \sum u_n$
- Max upload rate (limiting max download rate) =  $u_s + \sum u_n$
- Time to download the files  $n$  times by using the upload speed of all clients =  $nf/(u_s + \sum u_n)$

Time to distribute file ' $f$ ' to ' $n$ ' clients using peer-to-peer approach

- $T_{c-s} \geq \max \{f/u_s, f/d_{\min}, nf/(u_s + \sum u_n)\}$



# Client-Server vs P2P






# P2P File Distribution: BitTorrent

- File divided into 256Kb (it can be any size!) chunks
- Peers in torrent send/receive file chunks


*tracker*: tracks peers participating in torrent

*torrent*: group of peers exchanging chunks of a file




# P2P File Distribution: BitTorrent

- A new peer joining torrent:
  - has no chunks, but will accumulate them over time from other peers
  - registers with tracker to get list of peers and connects to a subset of peers ("neighbors") found on the tracker list.
- While downloading, a peer also uploads chunks to other peers
- A peer may change its connected peers with whom it exchanges chunks
- **Churn:** connected peers may come and go
- Once a peer has the entire file, it may (selfishly; **a leacher**) leave or (altruistically; **a seeder**) remain in the torrent (sharing its chunks with others)



# P2P File Distribution: BitTorrent

- (1) Alice "optimistically unchokes" Bob
- (2) Alice becomes one of Bob's top-four providers; Bob reciprocates
- (3) Bob becomes one of Alice's top-four providers



# P2P Example

Say, a torrent has 100 pieces/chunks.

Example 1:

- Client 1 has chunks 1 to 30 – Peer
- Client 2 has chunks 25 to 60 – Peer
- Client 3 has chunks 1 to 100 – Seeder
- Client 4 joins the torrent... Will this client be able to download the whole torrent?

Example 2:

- Client 1 has chunks 1 to 30 – Peer
- Client 2 has chunks 25 to 60 – Peer
- Client 3 has chunks 60 to 99 – Seeder
- Client 4 joins the torrent... Will this client be able to download the whole torrent?

# Content Distribution Networks (CDN)



**Content Delivery Network (CDN)**

# Video Streaming and CDNs: Context

- **Video traffic:** major consumer of Internet bandwidth
  - Netflix, YouTube: 37%, 16% of downstream residential ISP traffic
  - ~1B YouTube users, ~75M Netflix users
- **Challenges:**
  - **Scale:** how to reach ~1B users? single mega-video server won't work (why?)
  - **Heterogeneity:** different users have different capabilities (e.g., wired versus mobile; bandwidth rich versus bandwidth poor)
- **Solution:** distributed, application-level infrastructure






# Multimedia Video

- **Video:** sequence of images displayed at constant rate
  - e.g., 24 images/sec
- **Digital image:** array of pixels
  - each pixel represented by bits
- **Coding:** use redundancy **within** and **between** images to **decrease number of bits** used to encode image
  - spatial (within image)
  - temporal (from one image to next)

*spatial coding example:* instead of sending  $N$  values of same color (all purple), send only two values: color value (*purple*) and number of repeated values ( $N$ )



frame  $i$

*temporal coding example:* instead of sending complete frame at  $i+1$ , send only differences from frame  $i$



frame  $i+1$


# Multimedia: Video

- **CBR (constant bit rate):** video encoding rate fixed
- **VBR (variable bit rate):** video encoding rate changes as amount of spatial, temporal coding changes
- Examples:
  - MPEG 1 (CD-ROM) 1.5 Mbps
  - MPEG2 (DVD) 3-6 Mbps
  - MPEG4 (often used in Internet, < 1 Mbps)



# Streaming Stored Video

- Simple scenario



# Streaming: DASH

- **DASH: Dynamic, Adaptive Streaming over HTTP**
- **Server:**
  - divides video file into multiple chunks
  - each chunk stored, encoded at different rates
  - manifest file: provides URLs for different chunks
- **Client:**
  - periodically measures server-to-client bandwidth
  - consulting manifest, requests one chunk at a time
    - chooses maximum coding rate sustainable given current bandwidth
    - can choose different coding rates at different points in time (depending on available bandwidth at time)

# Streaming: DASH

- **“Intelligence” at client:** client determines
  - **when** to request chunk (so that **buffer** starvation, or overflow does not occur)
  - **what encoding rate** to request (higher quality when more bandwidth available)
  - **where** to request chunk (can request from URL server that is **“close” to client** or has **high available bandwidth**)


# Content Distribution Networks (CDNs)

**Challenge:** how to stream content (selected from millions of videos) to hundreds of thousands of simultaneous users?

- **Option 1:** single and large “mega-server”
  - Single point of failure
  - Long path to distant clients
  - Multiple copies of video sent over outgoing link
  - Single point of network congestion
  - This solution doesn’t scale
- **Option 2:** store multiple copies of videos at multiple distributed sites (CDN)
  - **Enter Deep:** push CDN servers deep into many access networks
    - close to users
    - used by Akamai, 1700 locations
  - **Bring Home:** smaller number (10’s) of larger clusters in **POPs near** (but not within) access networks
    - used by Limelight

# Content Distribution Networks (CDNs)

- CDN: stores copies of content at CDN nodes
  - e.g. Netflix stores copies of MadMen
- subscriber requests content from CDN
  - directed to nearby copy, retrieves content
  - may choose different copy if network path congested




# Content Distribution Networks (CDNs)

- **Over the top challenges:** coping with a congested Internet
  - From which CDN node to retrieve content for a user?
  - What's the viewer behavior in presence of a congestion?
  - What content to place in which CDN node?

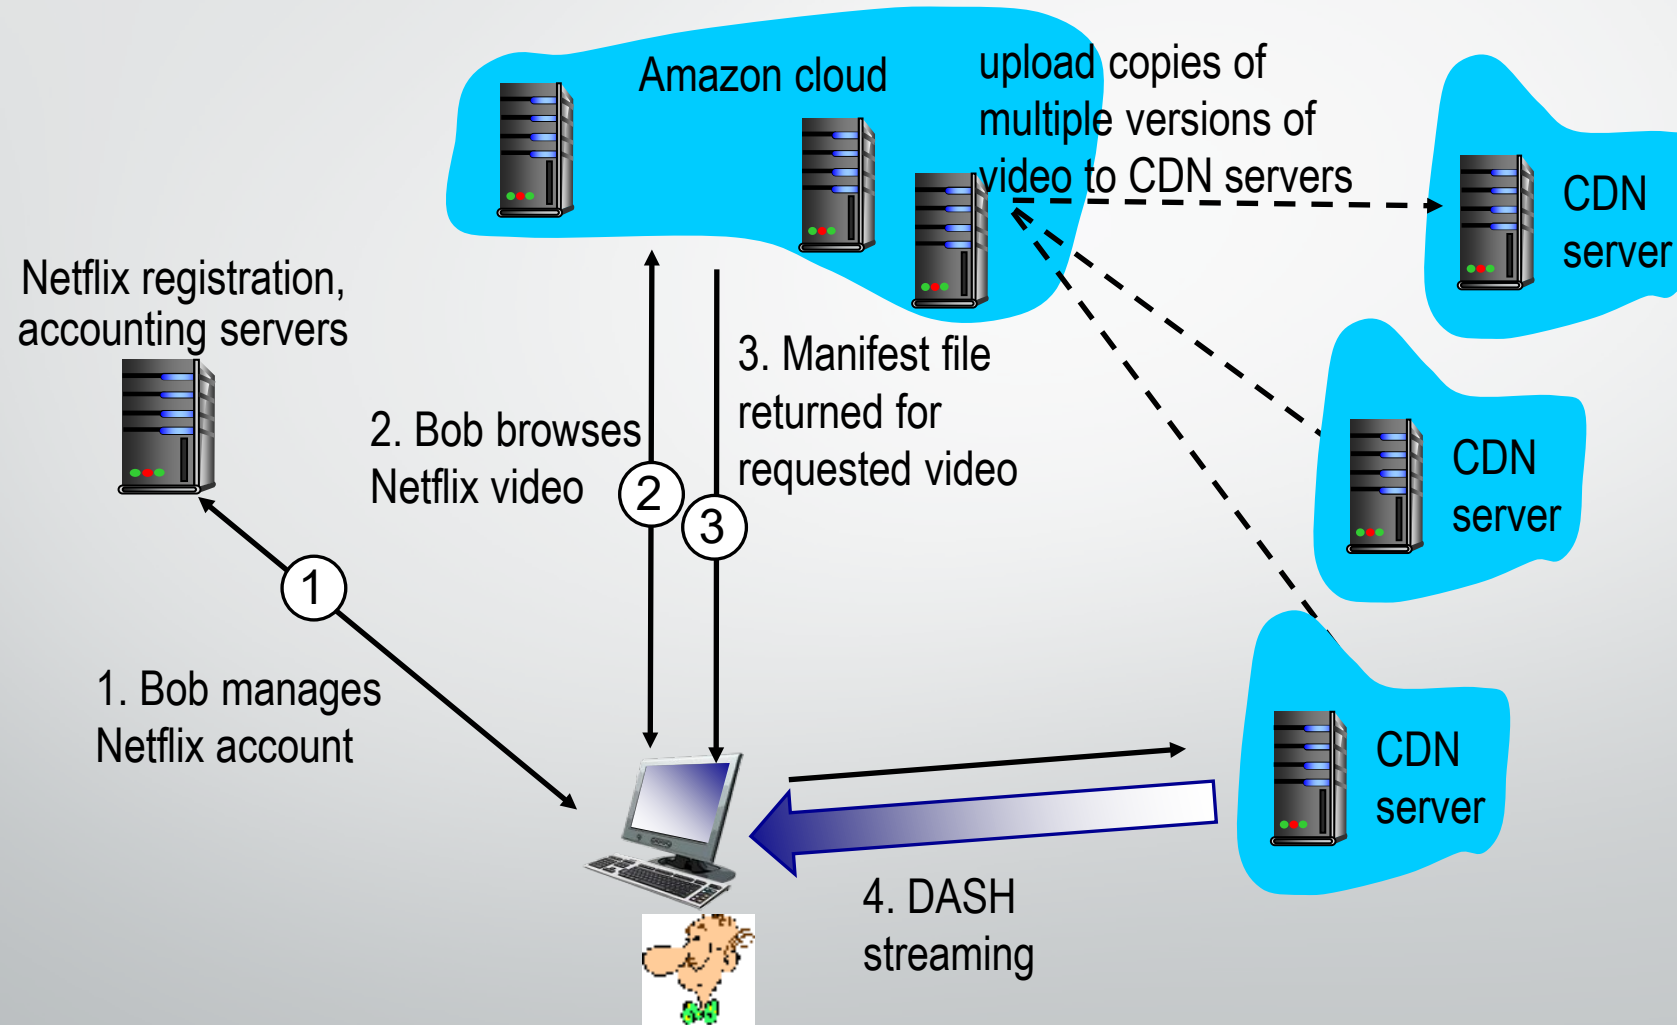
# CDN Content Access: A Closer Look

- Bob (client) requests video <http://netcinema.com/6Y7B23V>
  - video stored in CDN at <http://KingCDN.com/NetC6y&B23V>





# Case Study: Netflix





# The End

- **References**

- [1] Brownlee, M. [MKBHD]. (2019, October 12). This Is What Happens When You Re-Upload a YouTube Video 1000 Times! . Retrieved from <https://www.youtube.com/watch?v=JR4KHfqw-oE>
- [2] Kurose, J. F., & Ross, K. W. (2017). Computer networking: A top-down approach.