



Inspiring Excellence

# Introduction to Transport Layer

Lecture 6 | Part 1 | CSE421 – Computer Networks

Department of Computer Science and Engineering  
School of Data & Science

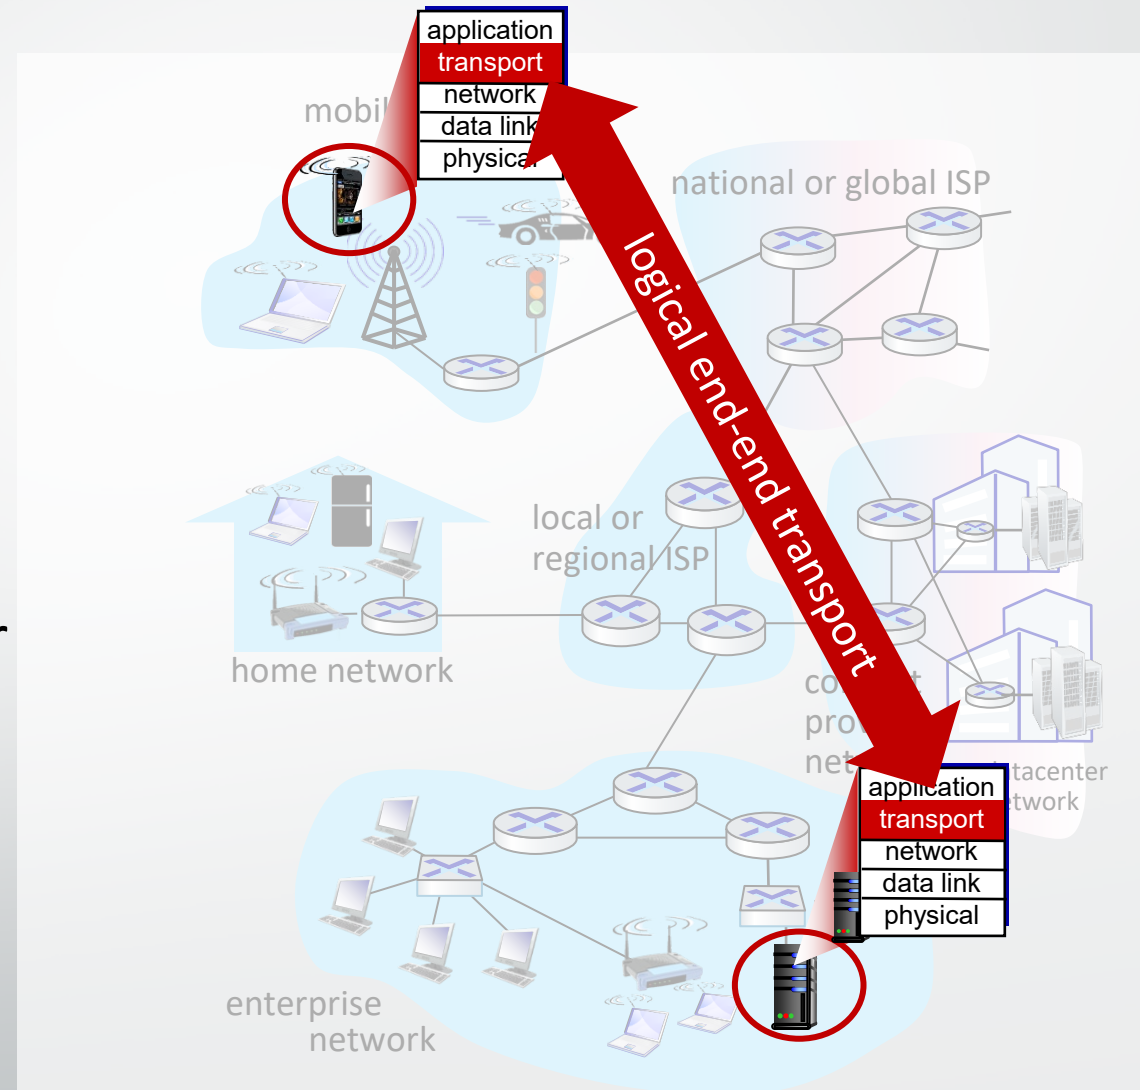
# Objectives

## our goals:

- understand principles behind transport layer services:
  - reliable data transfer, segmentation, flow control etc..
- learn about Internet transport layer protocols:
  - UDP: connectionless transport
  - TCP: connection-oriented reliable transport

# Transport services and protocols

- provide *logical communication* between application processes running on different hosts
- transport protocols actions in end systems:
  - sender: breaks application messages into *segments*, passes to network layer
  - receiver: reassembles segments into messages, passes to application layer
- two transport protocols available to Internet applications
  - TCP, UDP



# Transport vs. Network layer

- **network layer:** logical communication between hosts
- **transport layer:** logical communication between processes
  - relies on, enhances, network layer services

## *household analogy:*

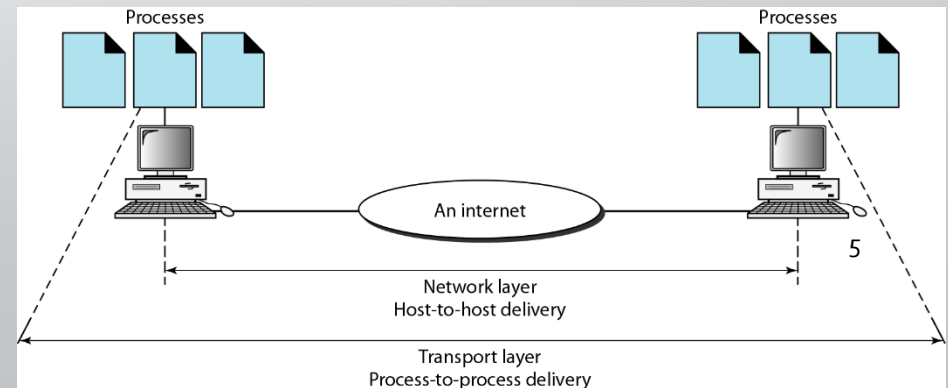
*12 kids in Ann's house sending letters to 12 kids in Bill's house:*

- hosts = houses
- processes = kids
- app messages = letters in envelopes
- transport protocol = Ann and Bill who demux to in-house siblings
- network-layer protocol = postal service

# Transport Layer

- The transport layer is responsible for the delivery of a message from one process (sender) to another (receiver).
- Transport Layer PDU is called **Segments**
- Functions:
  - Segmentation and Reassembly
  - Adds Port Address to identify the application
  - Multiplexing
  - Connection establishment and termination
  - Flow and Error Control

• \*PDU – Protocol Data Unit



# Purpose of the Transport Layer

- Primary responsibilities:
  - Segmenting the data and managing each piece.
  - Reassembling the segments into streams of application data.
  - Identifying the different applications.
  - Multiplexing
  - Initiating connection.
  - Performing end users.
  - Enabling error recovery.

Reliability

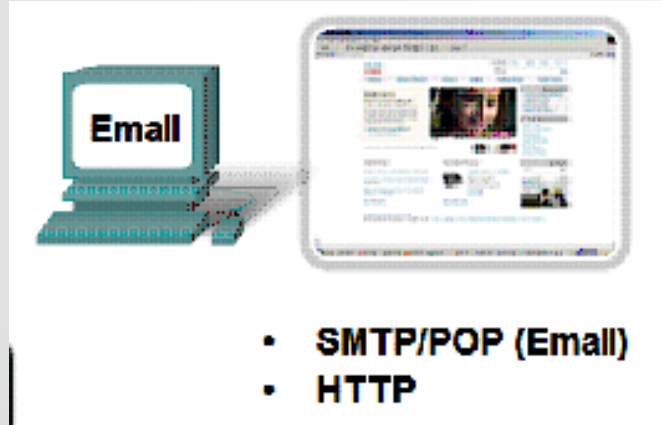
# Reliability

- Three **basic operations** of reliability are:
  - Initiating sessions and tracking transmitted data
  - Acknowledging received data
  - Retransmitting any unacknowledged data
- To support these reliability operations, more control data is added in the Layer 4 header.



# Different Applications

## Different Requirements

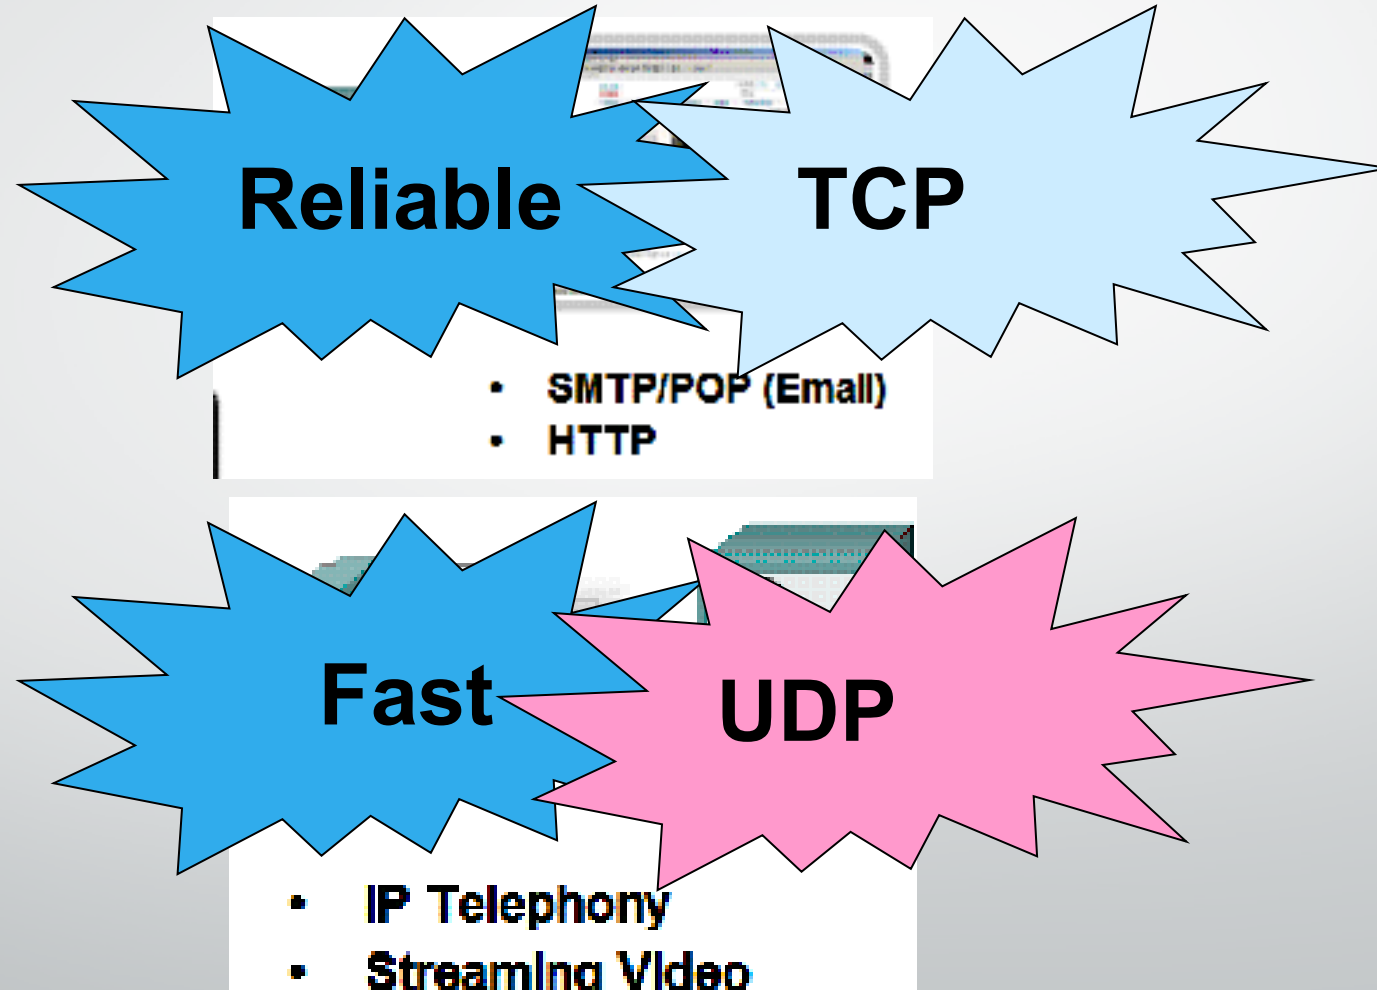


- Some applications need their data to be complete with no errors or gaps and they can accept a slight delay to ensure this.

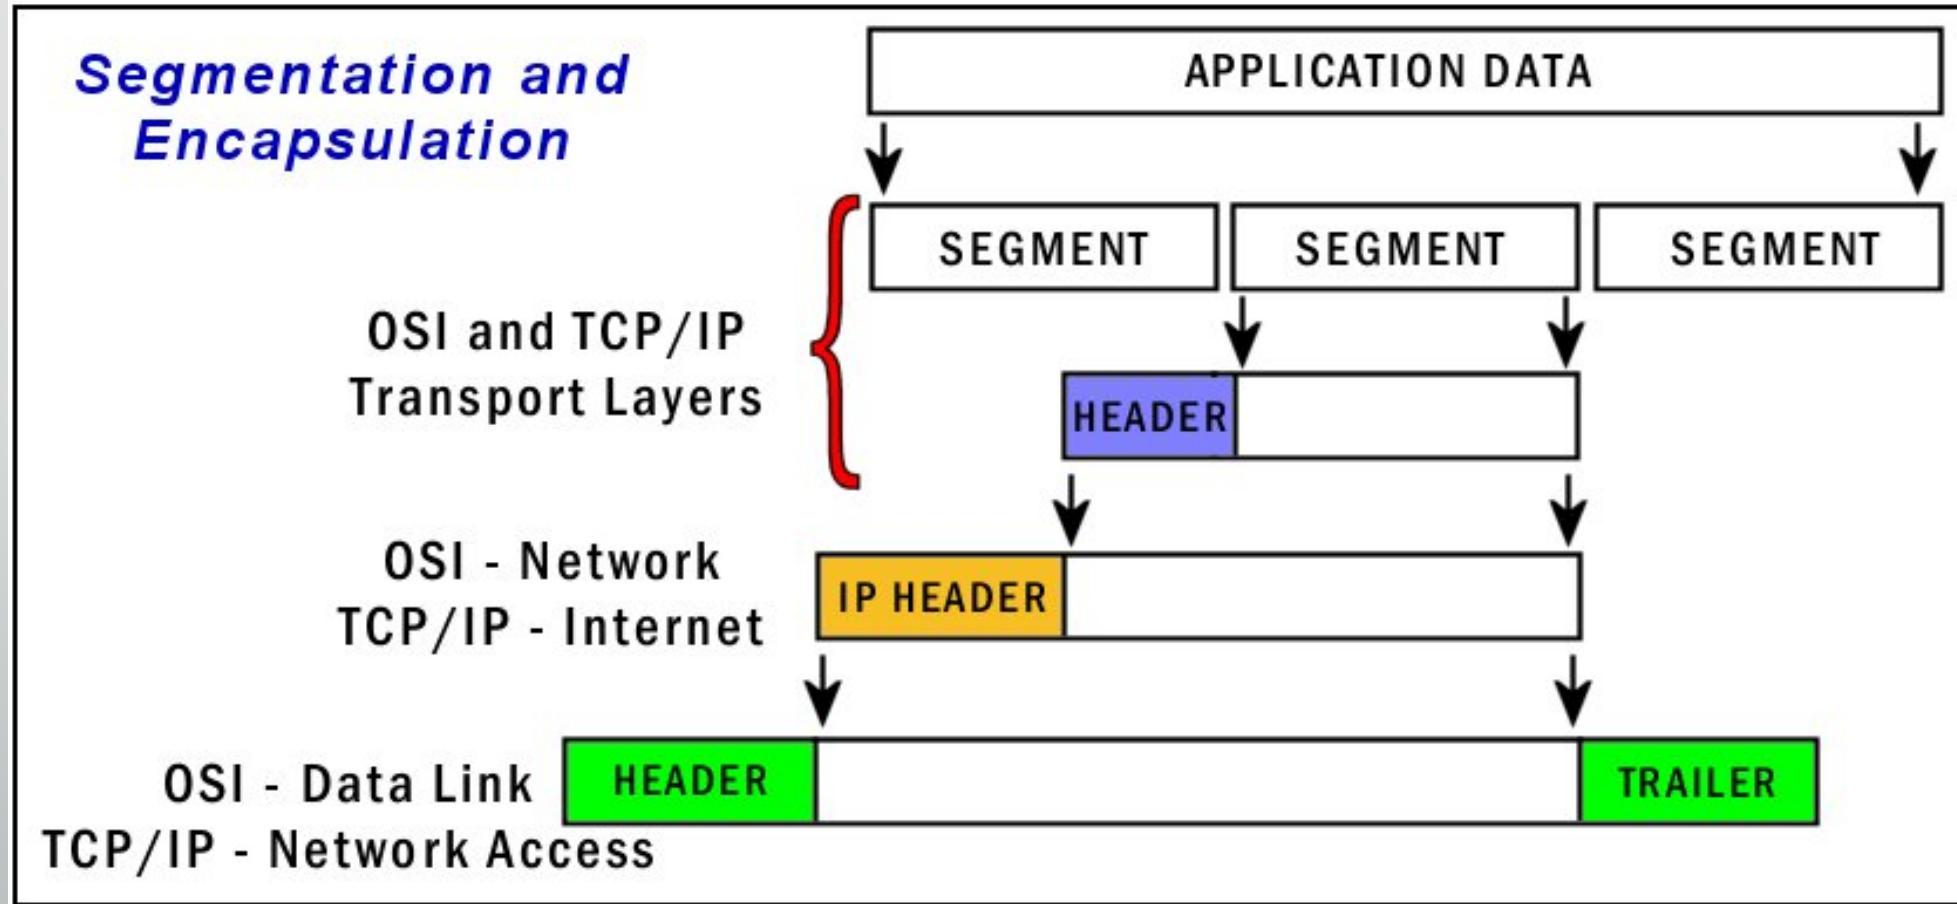
- Some applications can accept occasional errors or gaps in the data but they cannot accept any delay.



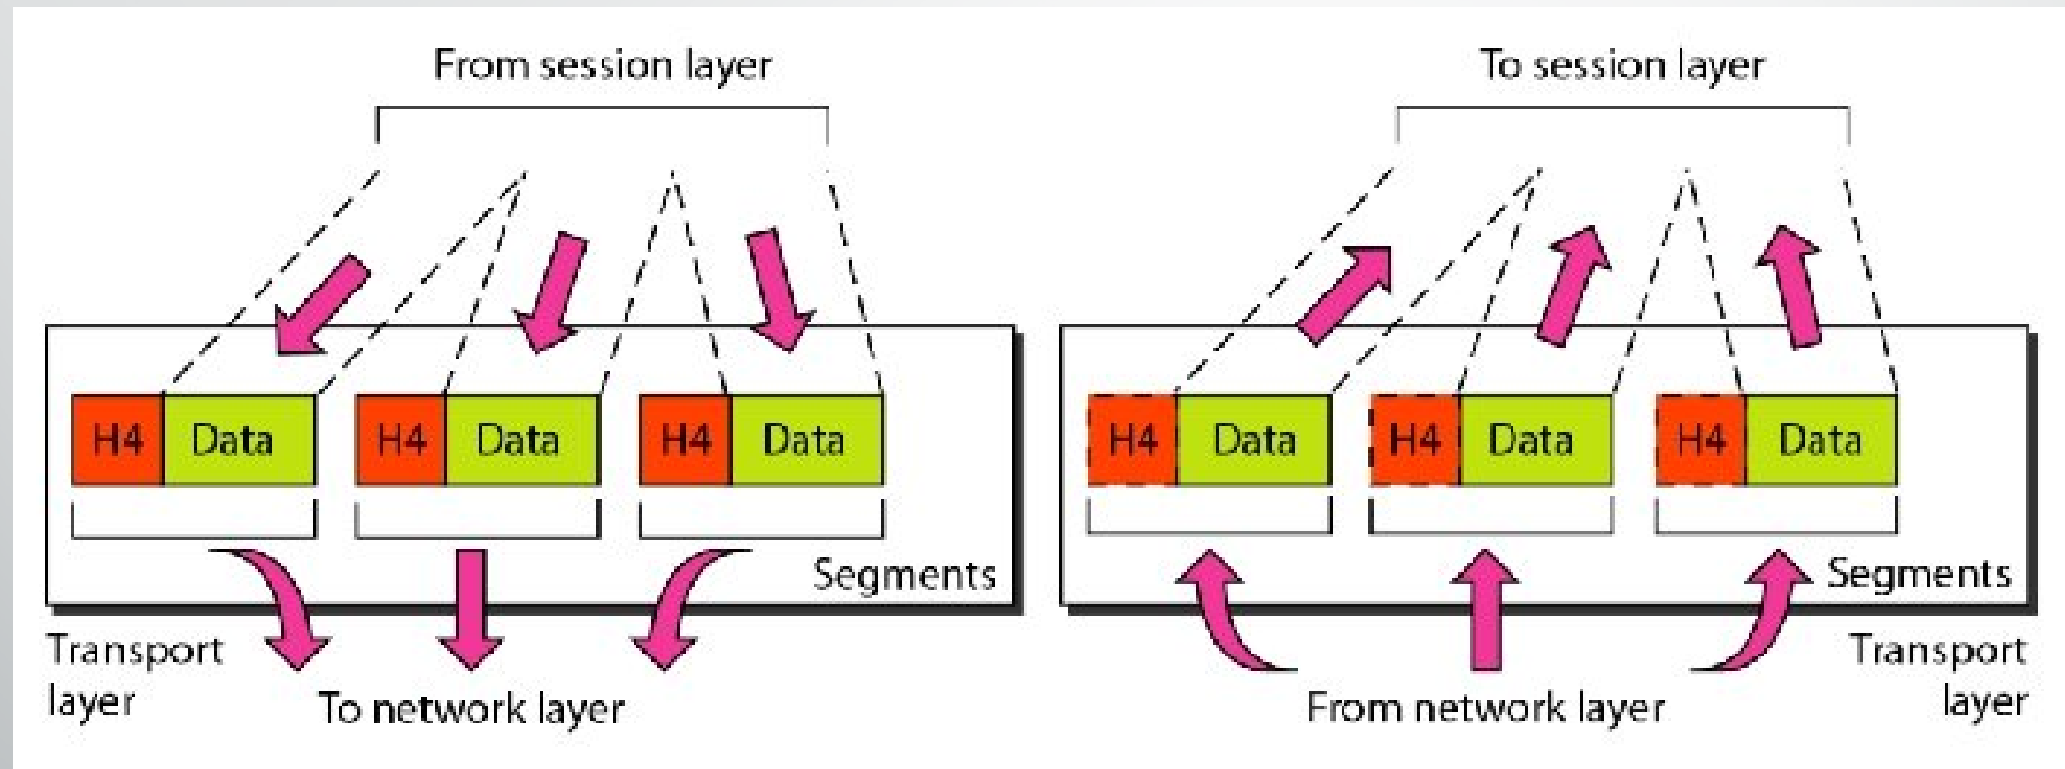
# Solution : Two transport protocols?



# Function 1 – Segmentation and Reassembly

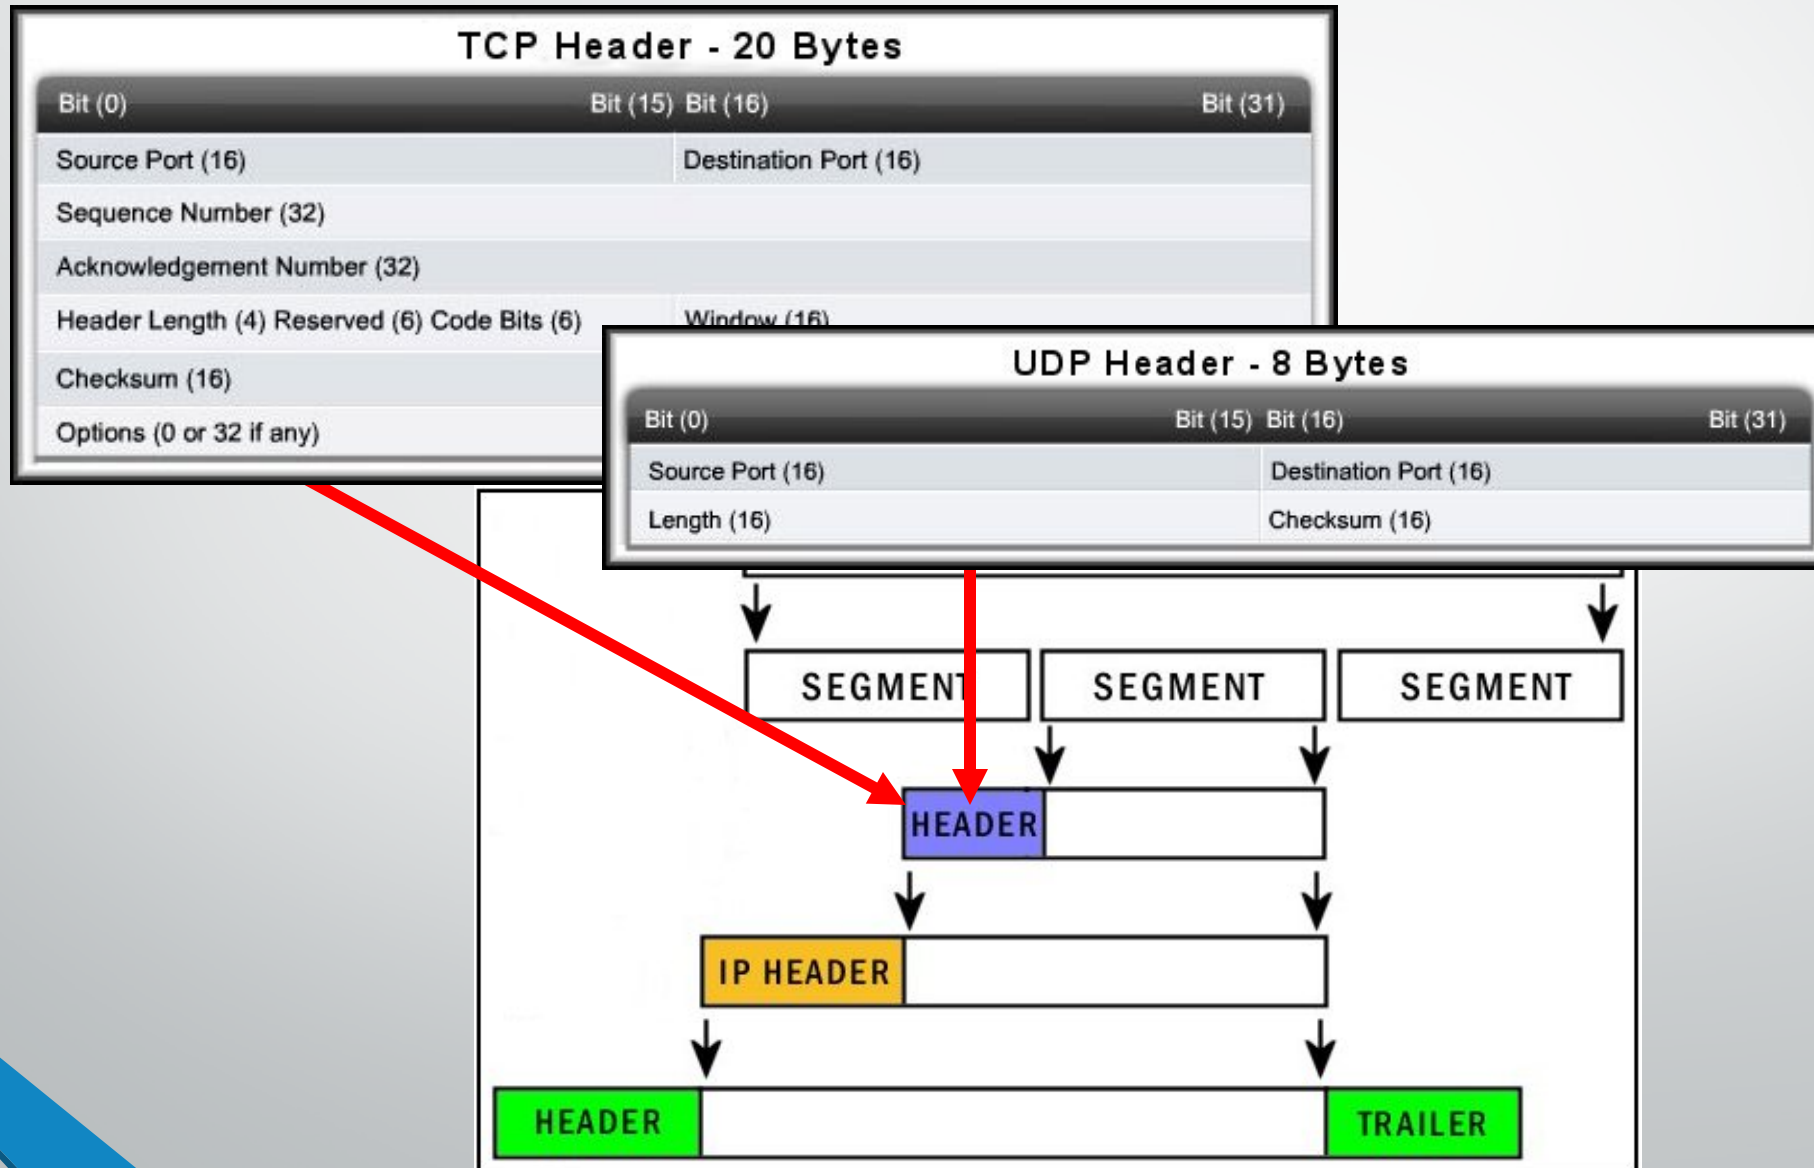


# Function 1 – Segmentation and Reassembly

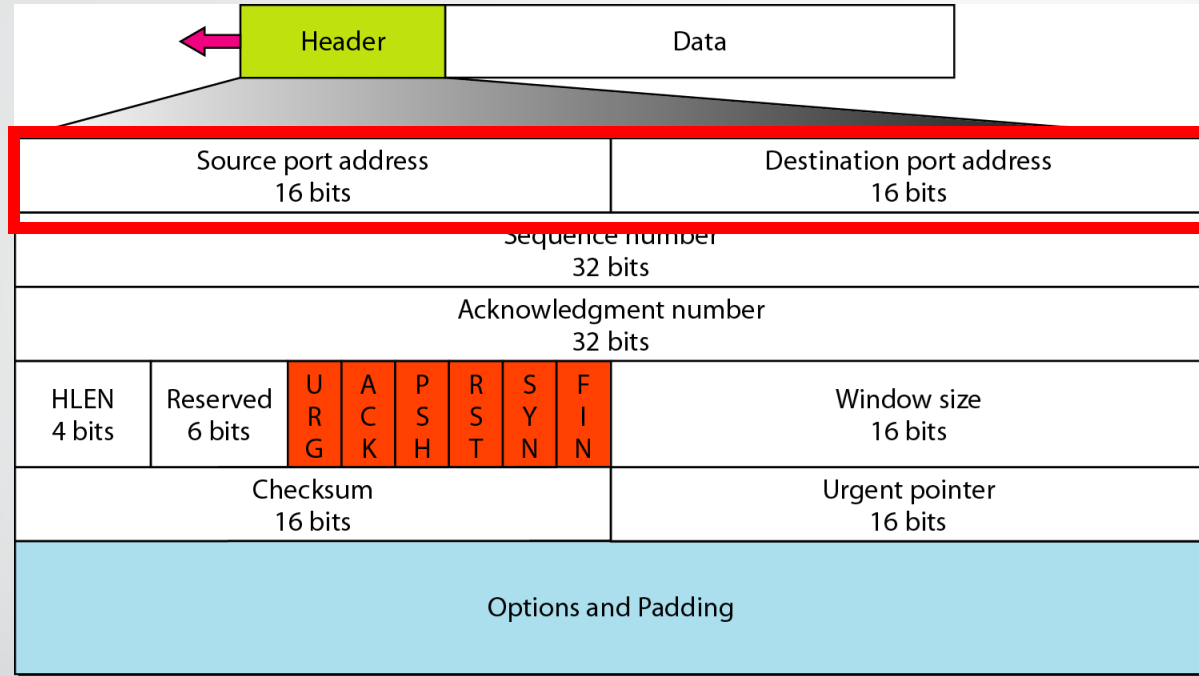


- Also known as encapsulation and de-capsulation.

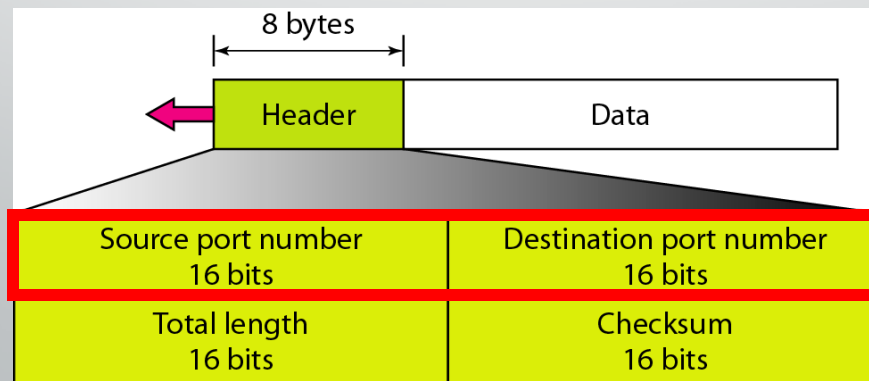
# TCP and UDP Headers



# TCP and UDP Headers



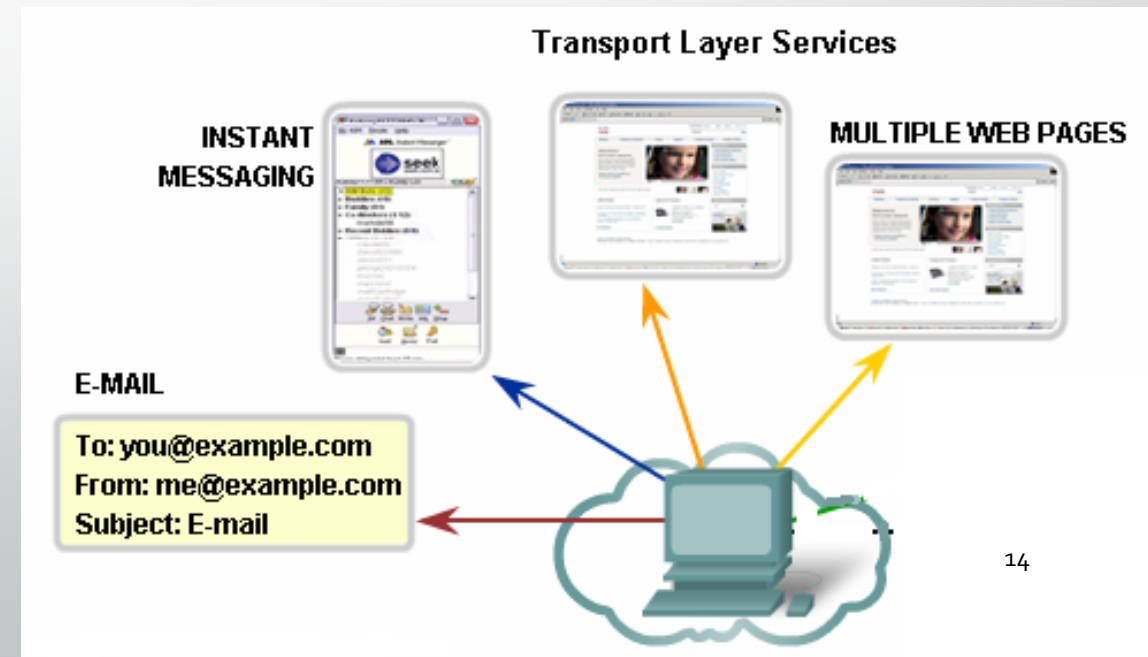
? TCP Header



? UDP Header


# Function2 – Identification Applications Using Port Address

- Port Numbers/Addresses are used to identify different applications/processes running in a computer
- 16-bit in length
  - Represented as one single decimal number
  - Range **0 - 65535**
  - e.g. **80 – Web**; **25 – SMTP**
  - **110 – POP3**, **531 – Instant Messaging**



# Port Numbers

- Internet Corporation for Assigned Names and Numbers (ICANN) assigns port numbers.
- **Three** categories:



Port Number Range	Port Group
0 to 1023	Well Known (Contact) Ports
1024 to 49151	Registered Ports
49152 to 65535	Private and/or Dynamic Ports



# Port Addressing Types

- Well-Known Ports:
  - Reserved for common services and applications.
  - Pre-assigned by ICANN

Port Number Range	Port Group
0 to 1023	Well Known (Contact) Ports
1024 to 49151	Registered Ports
49152 to 65535	Private and/or Dynamic Ports

20 – FTP Data

25 – SMTP

443 – HTTPS

21 – FTP Control

69 – TFTP

23 – Telnet

110 – POP3

520 – RIP

# Port Addressing Types

- Registered Ports:

- Port numbers that companies and other users register with ICANN for use by the applications that communicate using any one of the transport layer protocols.

Port Number Range	Port Group
0 to 1023	Well Known (Contact) Ports
1024 to 49151	Registered Ports
49152 to 65535	Private and/or Dynamic Ports

8008 – Alternate HTTP

1863 – MSN Messenger

8080 – Alternate HTTP

5004 – RTP

5060 – SIP (VoIP)

# Port Addressing Types

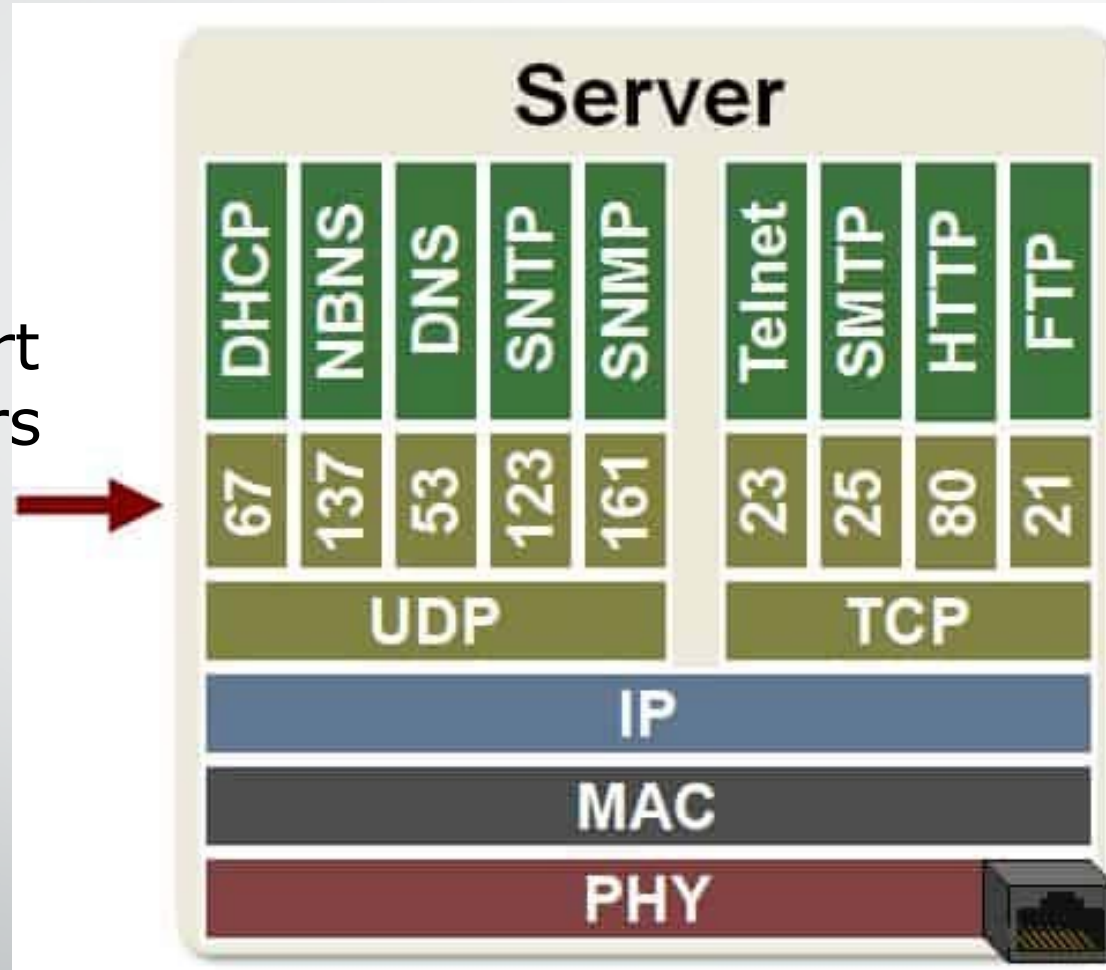
- Dynamic Ports:
  - Assigned to a user application at connect time.

Port Number Range	Port Group
0 to 1023	Well Known (Contact) Ports
1024 to 49151	Registered Ports
49152 to 65535	Private and/or Dynamic Ports

Dynamic port usage will become clearer as we move through the material.

# Port Numbers

- UDP Port Numbers






TCP Port Numbers

# Function 3 - Multiplexing

HTTP server


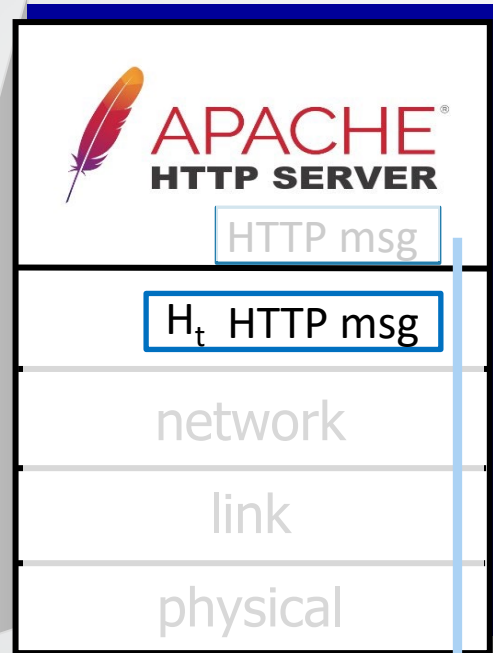


client



# Multiplexing

HTTP server


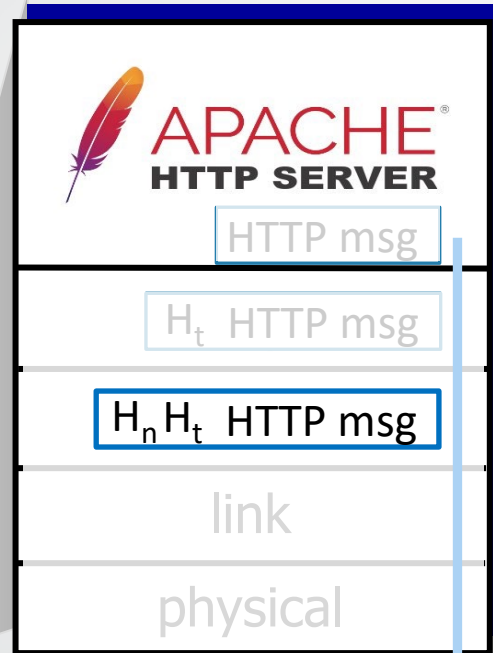


client



# Multiplexing

HTTP server

client

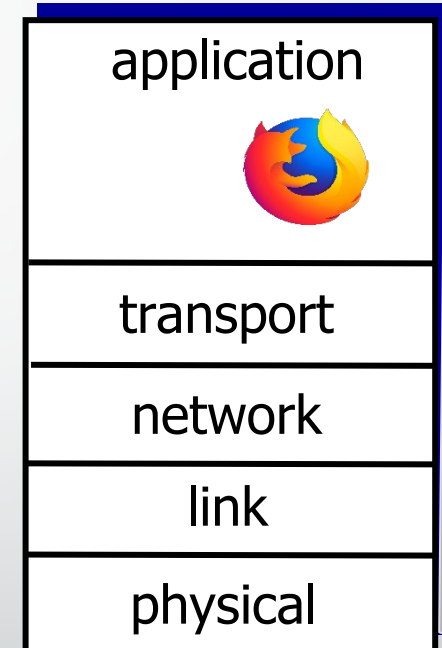

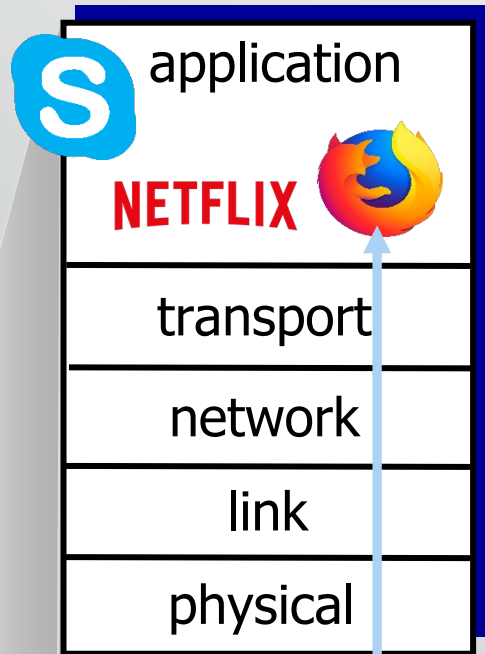




# Multiplexing

HTTP server

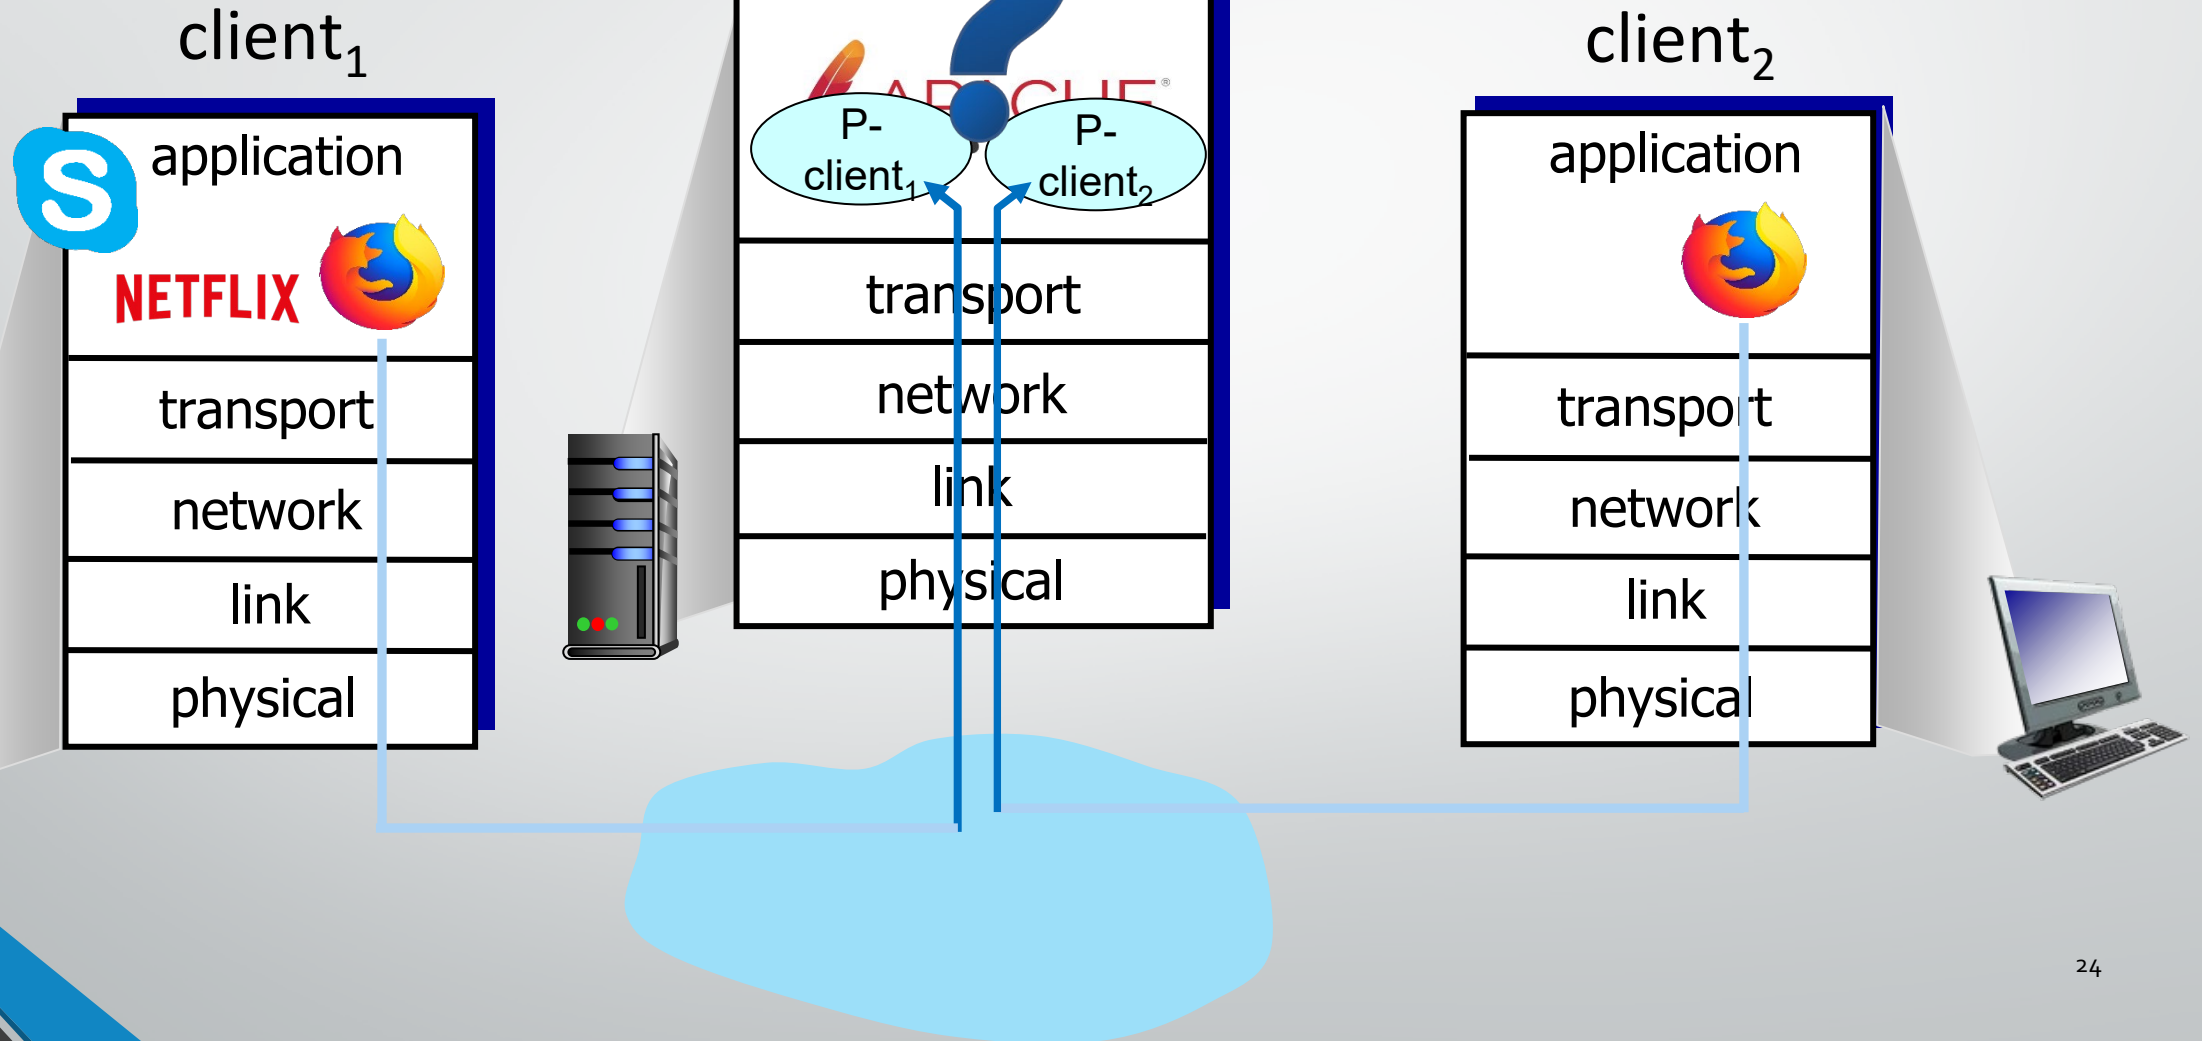
client



$H_n H_t$  HTTP msg

# Multiplexing

HTTP server



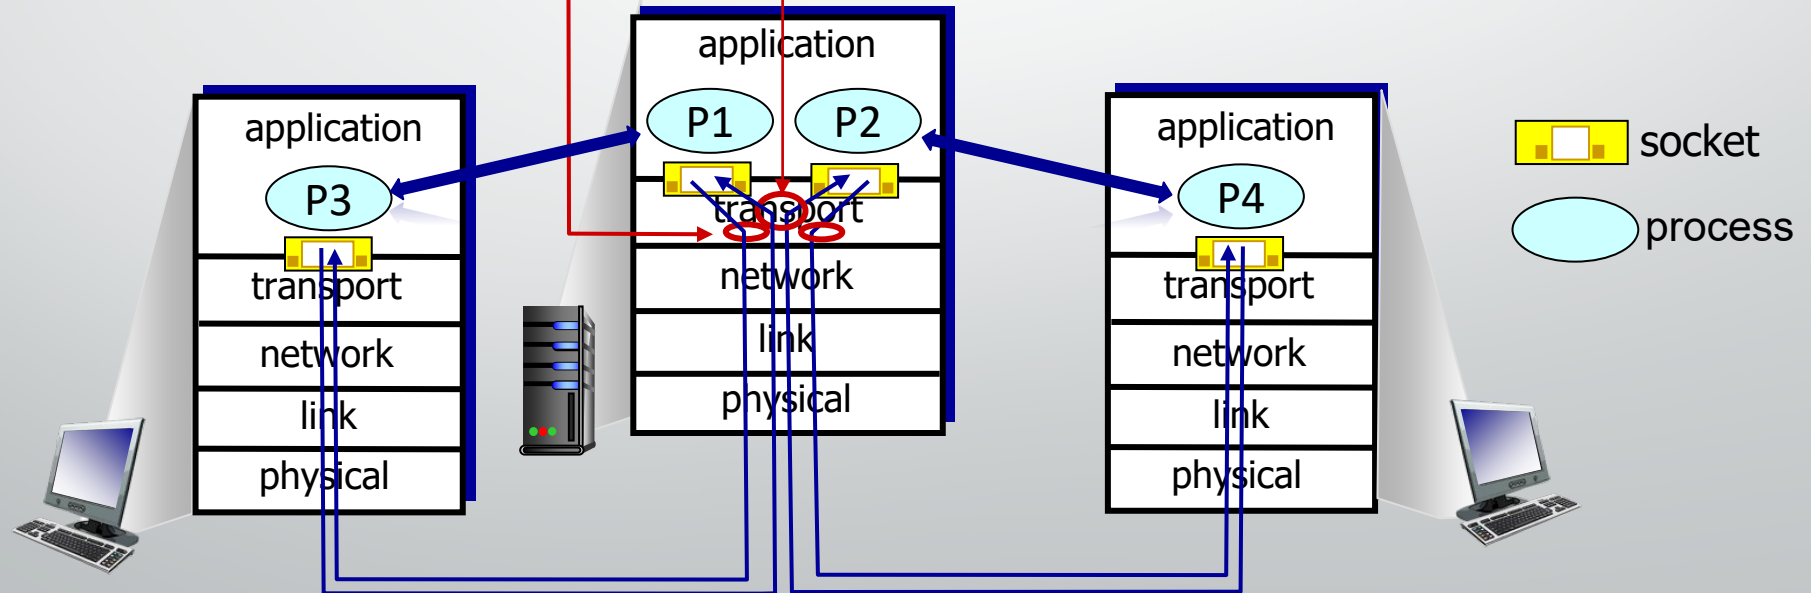
# Multiplexing/demultiplexing

## *multiplexing at sender:*

handle data from multiple sockets, add transport header (later used for demultiplexing)

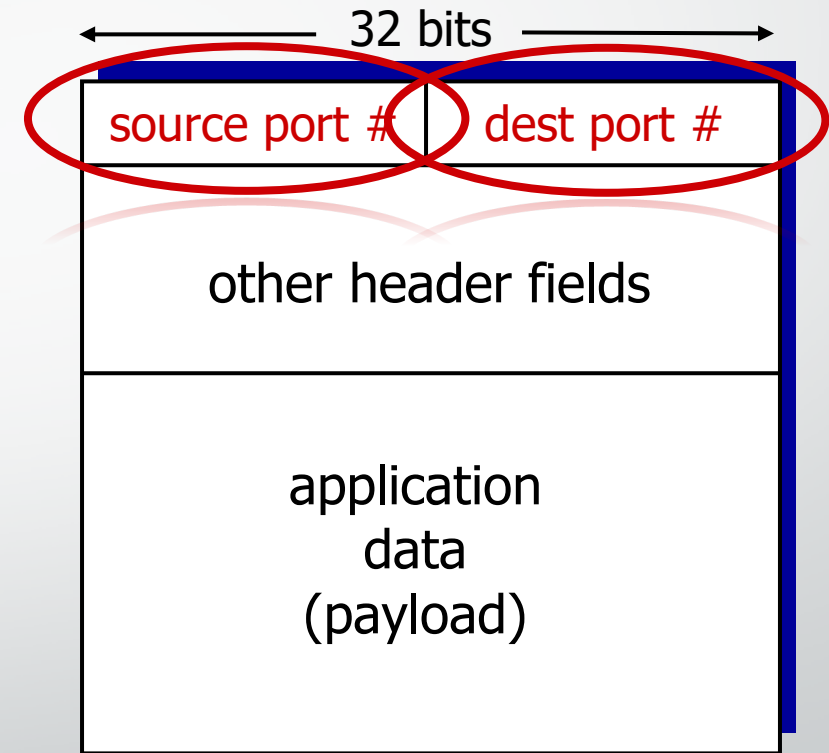
## *demultiplexing at receiver:*

use header info to deliver received segments to correct socket



# How demultiplexing works

- host receives IP datagrams
  - each datagram has source IP address, destination IP address
  - each datagram carries one transport-layer segment
  - each segment has source, destination port number
- host uses *IP addresses & port numbers* to direct segment to appropriate socket



TCP/UDP segment format

# Connectionless demultiplexing

- when creating datagram to send into UDP socket, must specify
  - destination IP address
  - destination port #

when receiving host receives *UDP* segment:

- checks destination port # in segment
- directs UDP segment to socket with that port #



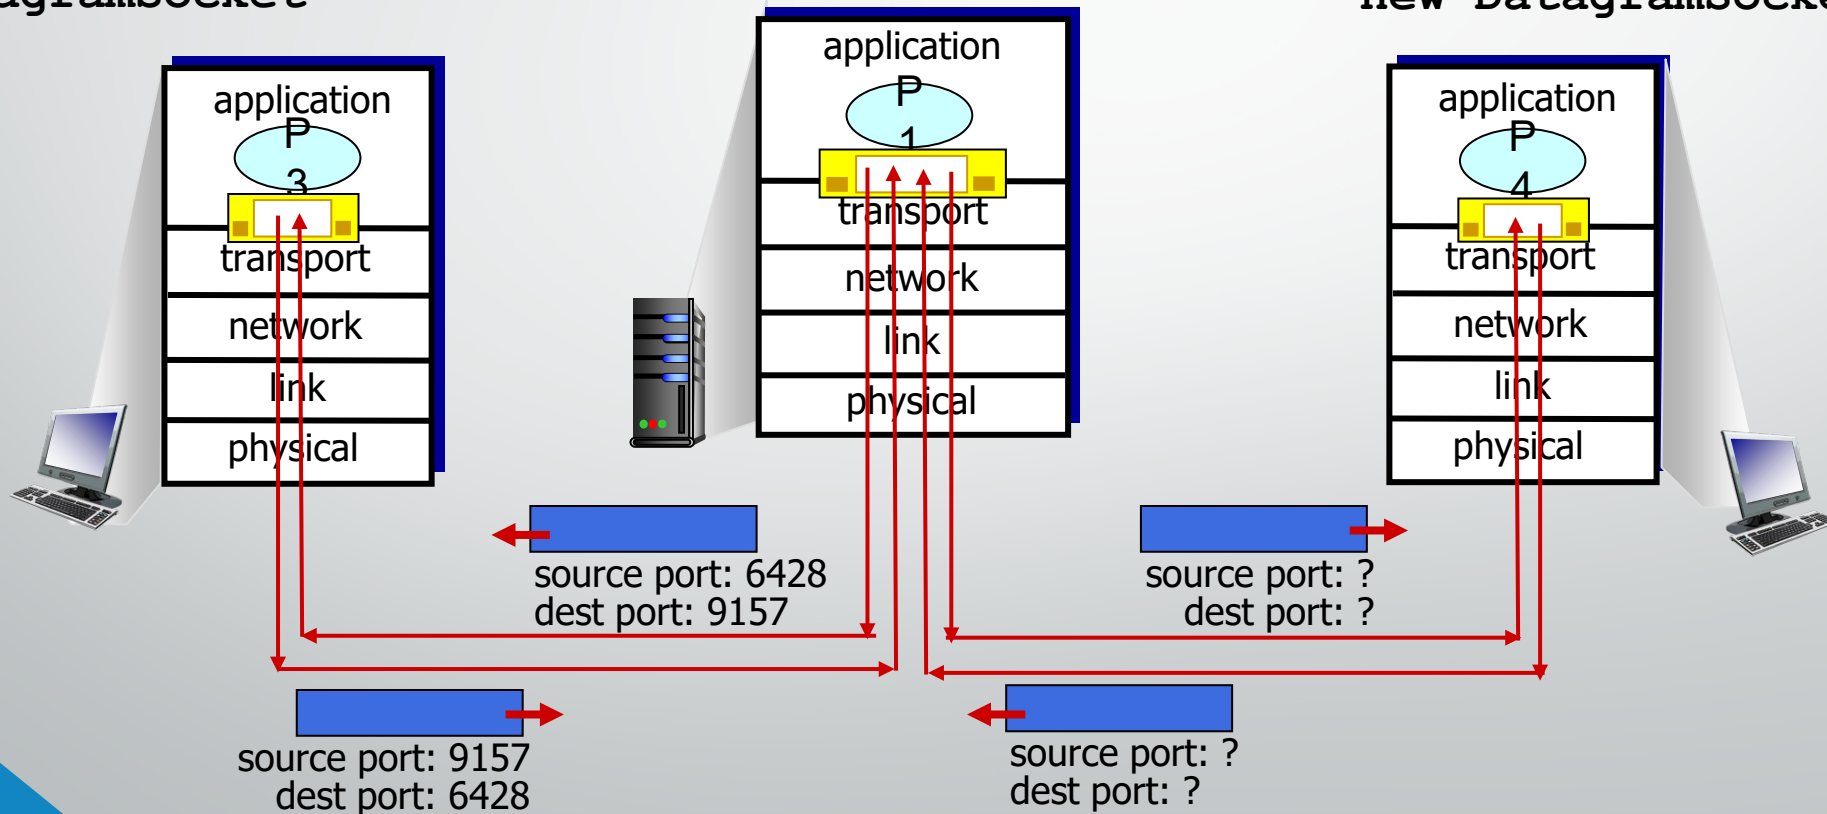
IP/UDP datagrams with *same dest. port #*, but different source IP addresses and/or source port numbers will be directed to *same socket* at receiving host

# Connectionless demultiplexing: an example

```
DatagramSocket mySocket2 =  
new DatagramSocket  
(9157);
```

```
DatagramSocket  
serverSocket = new  
DatagramSocket  
(6428);
```

```
DatagramSocket mySocket1 =  
new DatagramSocket (5775);
```

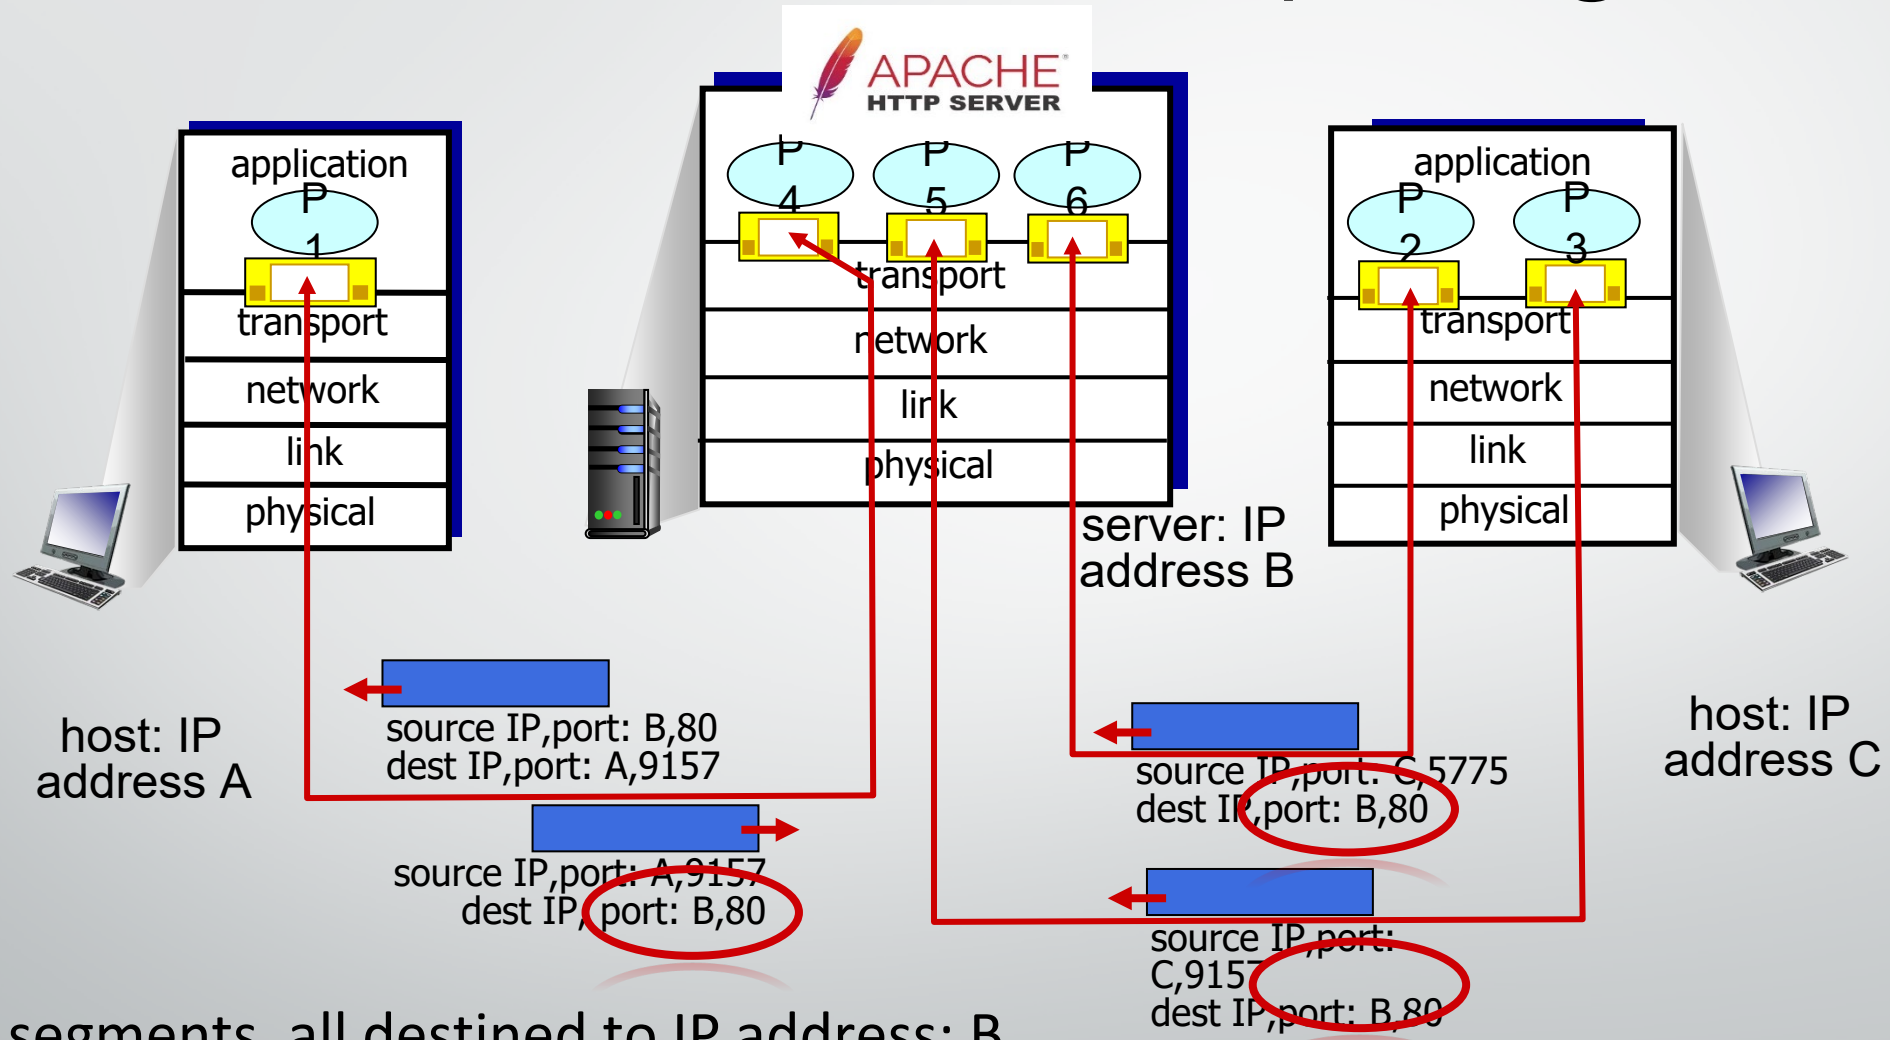


# Connection-oriented demultiplexing

- TCP socket identified by 4-tuple:
  - source IP address
  - source port number
  - dest IP address
  - dest port number
- demux: receiver uses *all four values (4-tuple)* to direct segment to appropriate socket
- server may support many simultaneous TCP sockets:
  - each socket identified by its own 4-tuple
  - each socket associated with a different connecting client



# Connection-oriented demultiplexing: example

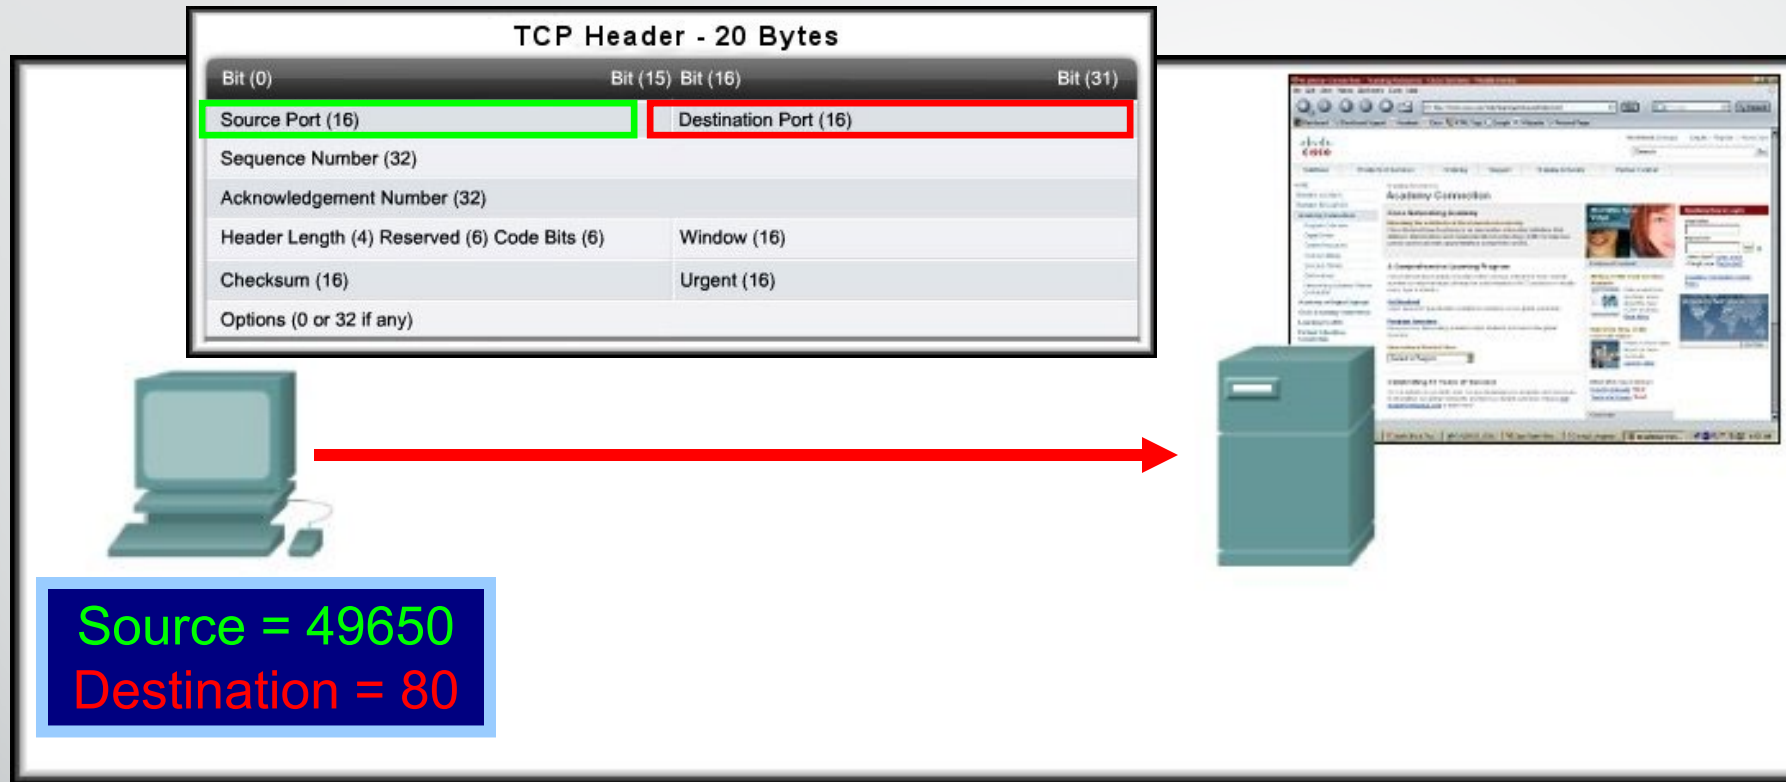


Three segments, all destined to IP address: B,  
dest port: 80 are demultiplexed to *different* sockets

# Summary

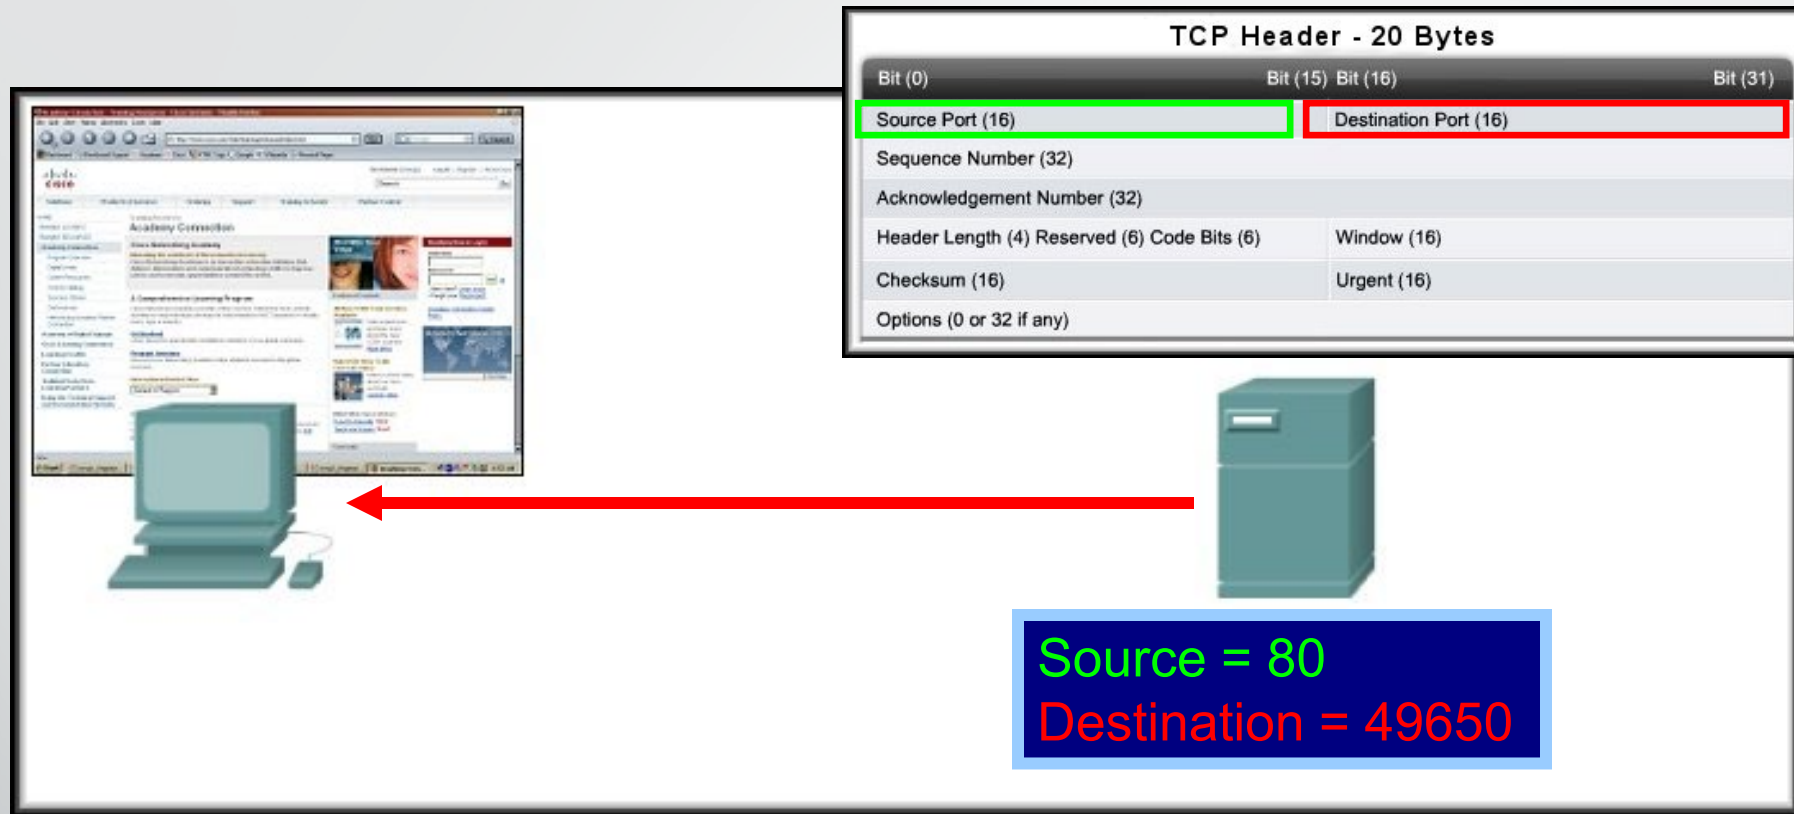
- Multiplexing, demultiplexing: based on segment, datagram header field values
- **UDP:** demultiplexing using destination port number (only)
- **TCP:** demultiplexing using 4-tuple: source and destination IP addresses, and port numbers
- Multiplexing/demultiplexing happen at *all* layers

# More on Port Numbers



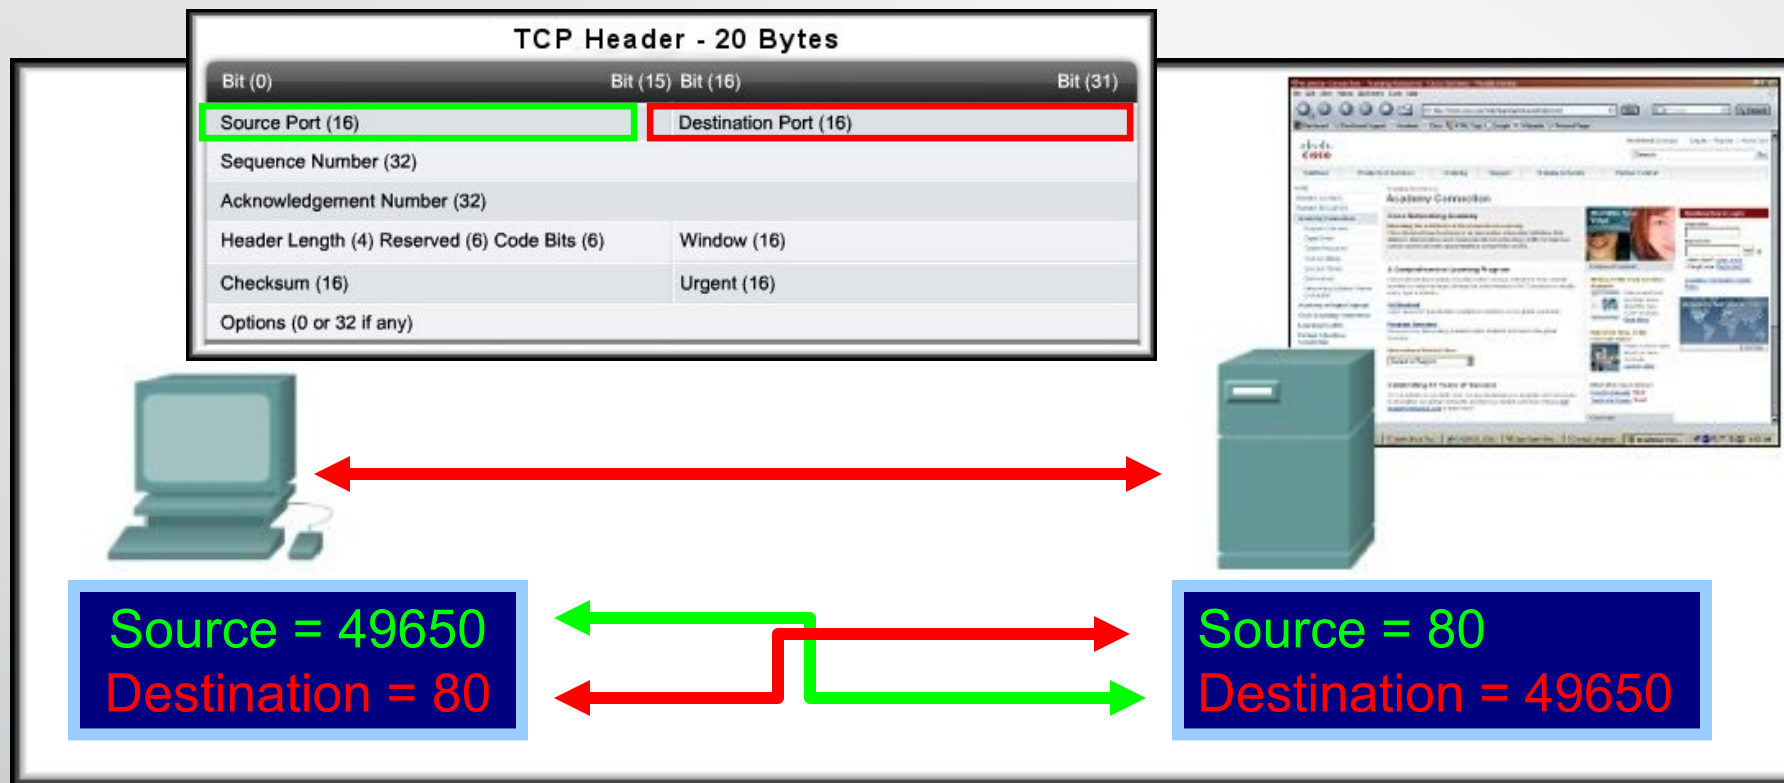
- Server is listening on Port 80 for HTTP connections.
- The client sets the destination port to 80 and uses a dynamic port as its source.

# Port Numbers in Action



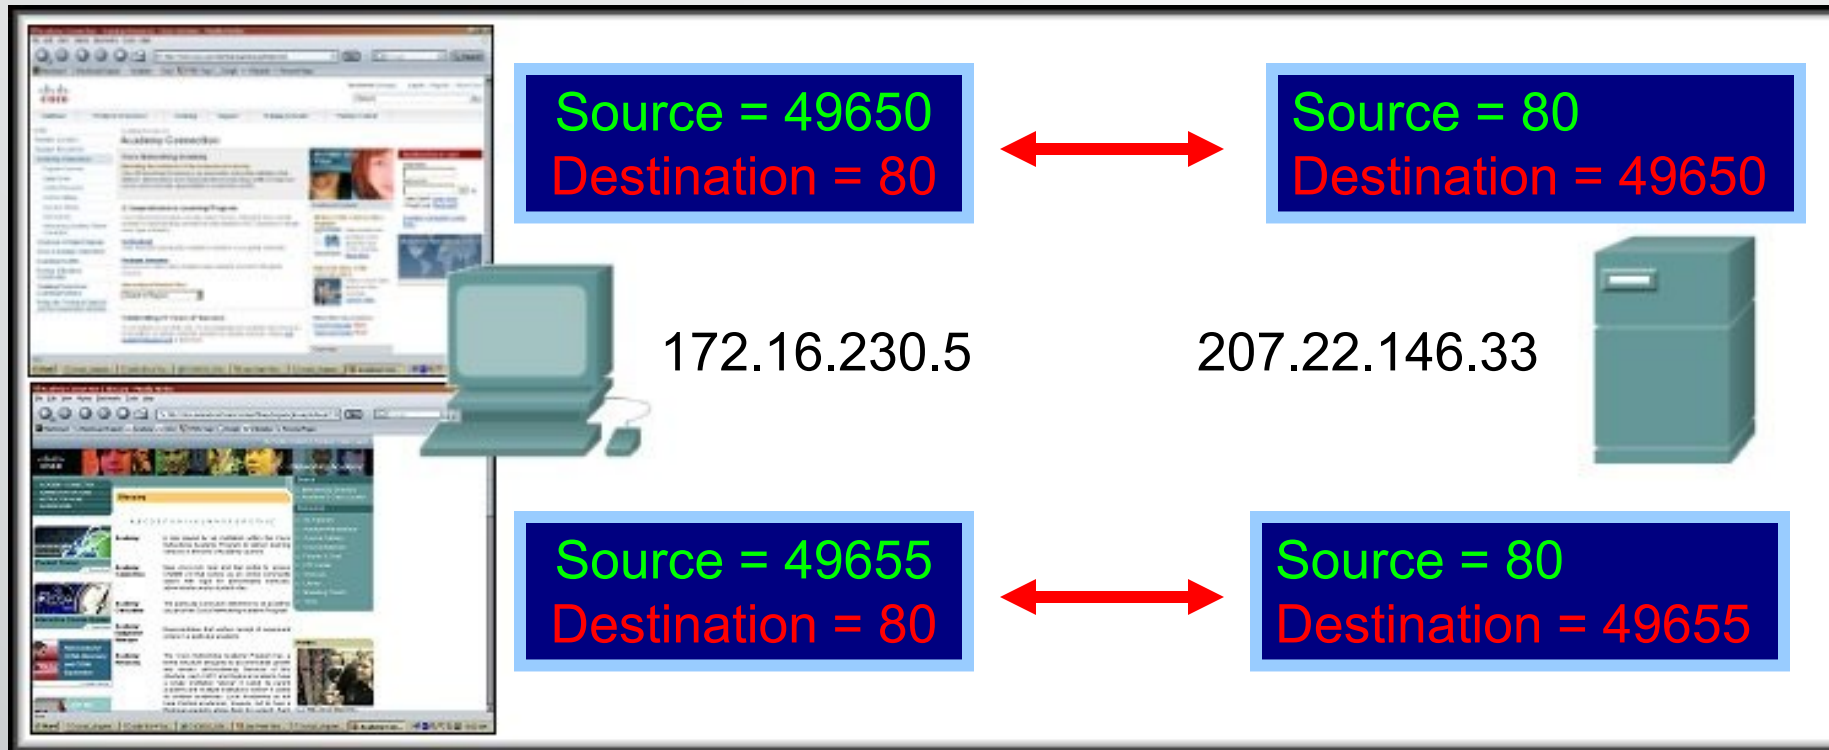
- Server replies with the web page.
  - Sets the source port to 80 and uses the client's source port as the destination.

# Port Numbers in Action



- Notice how the source and destination ports are used.
- Clients can use any random port number, servers can't.
  - Because clients won't be able to identify server process otherwise
- Servers, however, cannot use any random port number
  - Use of well-known port numbers!

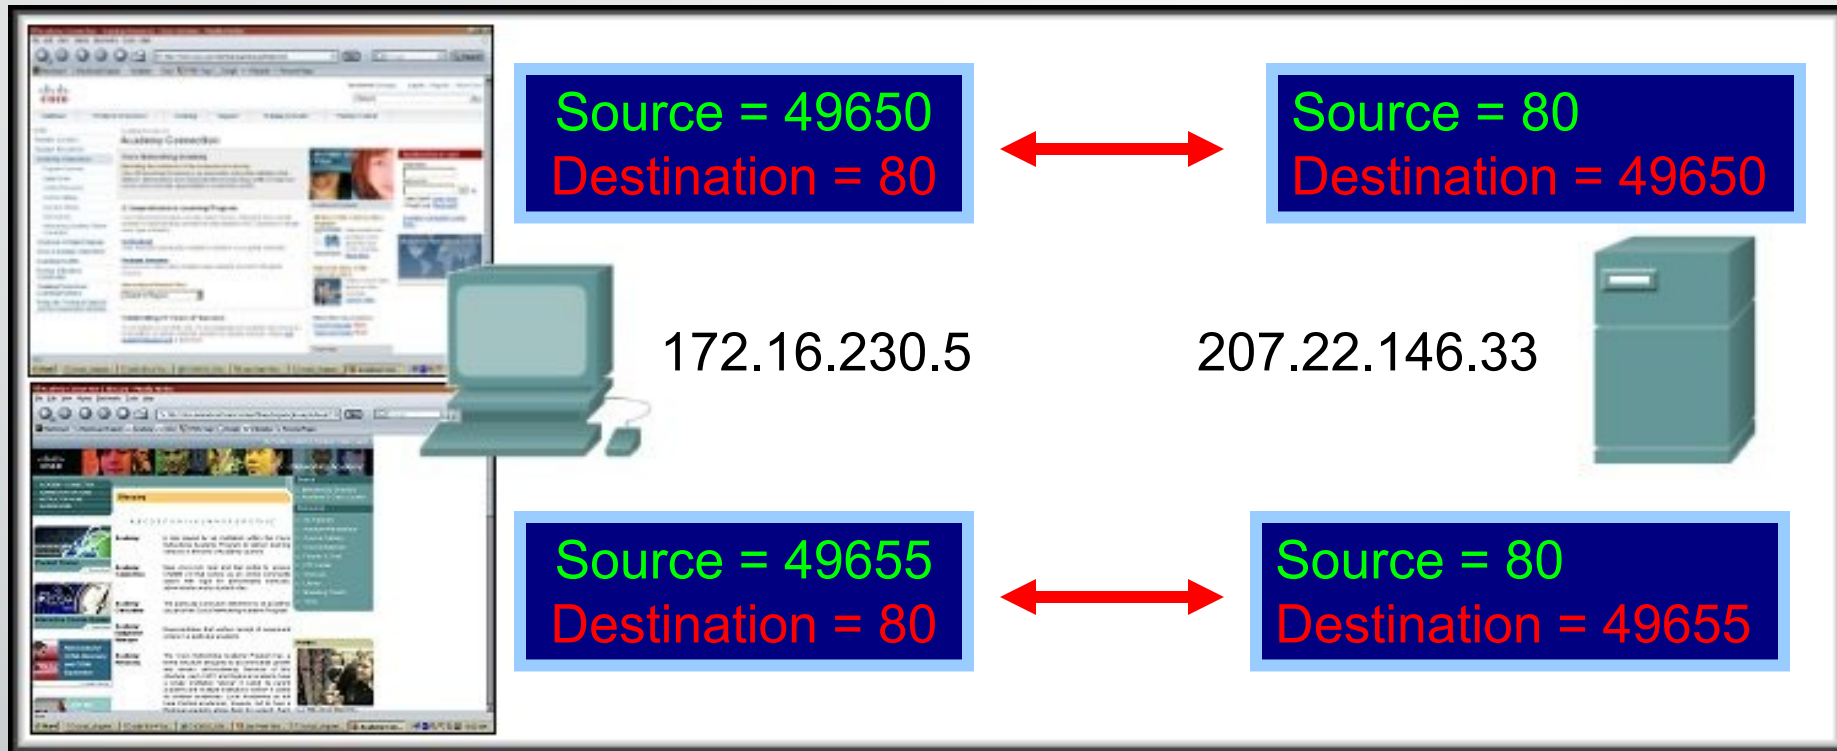
# Port Numbers in Action



- What if there are two sessions to the same server?
  - The client uses **another dynamic port** as its source and the destination is **still port 80**.
  - **Different source ports** keep the sessions unique on the server.



# Port Numbers in Action



- How does the Transport Layer keep them separate?
  - The socket (IP Address:Port)

172.16.230.5:49650 ↔ 207.22.146.33:49650  
172.16.230.5:49655 ↔ 207.22.146.33:49655



```
C:\WINDOWS\system32\cmd.exe
C:\>netstat -a -n

Active Connections

Proto Local Address          Foreign Address         State
TCP 0.0.0.0:135             0.0.0.0:*               LISTENING
TCP 0.0.0.0:445             0.0.0.0:*               LISTENING
TCP 0.0.0.0:3389            0.0.0.0:*               LISTENING
TCP 0.0.0.0:48698          0.0.0.0:*               LISTENING
TCP 127.0.0.1:3582         127.0.0.1:3583          ESTABLISHED
TCP 127.0.0.1:3583         127.0.0.1:3582          ESTABLISHED
TCP 192.168.1.103:135     204.225.7.4:80          TIME_WAIT
TCP 192.168.1.103:3586   0.0.0.0:*               LISTENING
UDP 0.0.0.0:445             *:.*                    *:*
UDP 0.0.0.0:500            *:.*                    *:*
UDP 0.0.0.0:1025           *:.*                    *:*
UDP 0.0.0.0:1026           *:.*                    *:*
UDP 0.0.0.0:1030           *:.*                    *:*
UDP 0.0.0.0:1038           *:.*                    *:*
UDP 0.0.0.0:1346           *:.*                    *:*
UDP 0.0.0.0:4500           *:.*                    *:*
UDP 0.0.0.1:123            *:.*                    *:*
UDP 127.0.0.1:1900          *:.*                    *:*
UDP 127.0.0.1:3492          *:.*                    *:*
UDP 192.168.1.103:123      *:.*                    *:*
UDP 192.168.1.103:137      *:.*                    *:*
UDP 192.168.1.103:138      *:.*                    *:*
UDP 192.168.1.103:139      *:.*                    *:*
netstat -a -n command
```

## Netstat - Network Utility Tool

- Actually, when you open up a single web page, there are usually several TCP sessions created, not just one.

# UDP: User Datagram Protocol [RFC 768]

*UDP is a connectionless, unreliable protocol that has **no flow and error control**. It uses port numbers to track and multiplex data received from the application layer.*

## why is there a UDP?

- ❖ no connection establishment (which can add delay)
- ❖ simple: no connection state at sender, receiver
- ❖ small header size
- ❖ no congestion control: UDP can blast away as fast as desired

# UDP

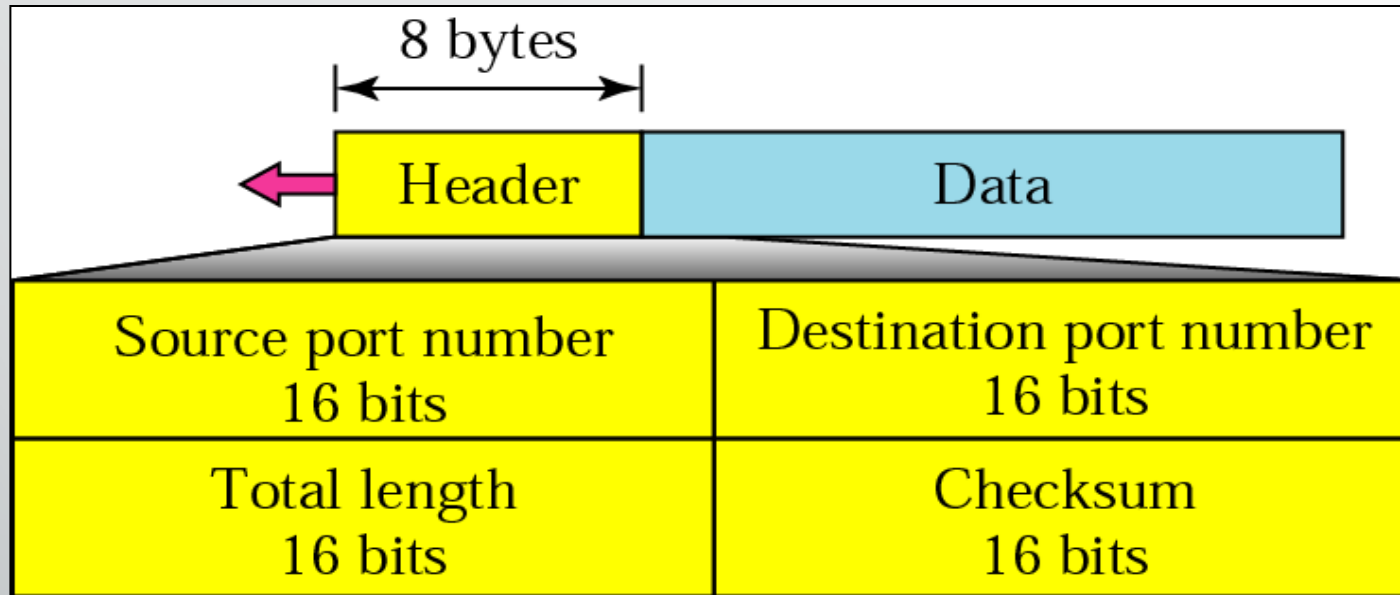
## ❖ UDP use:

- streaming multimedia apps (loss tolerant, rate sensitive)
- DNS
- SNMP
- HTTP/3

## ❖ reliable transfer over UDP:

- add reliability at application layer
- application-specific error recovery!

# User Datagram Protocol (UDP)

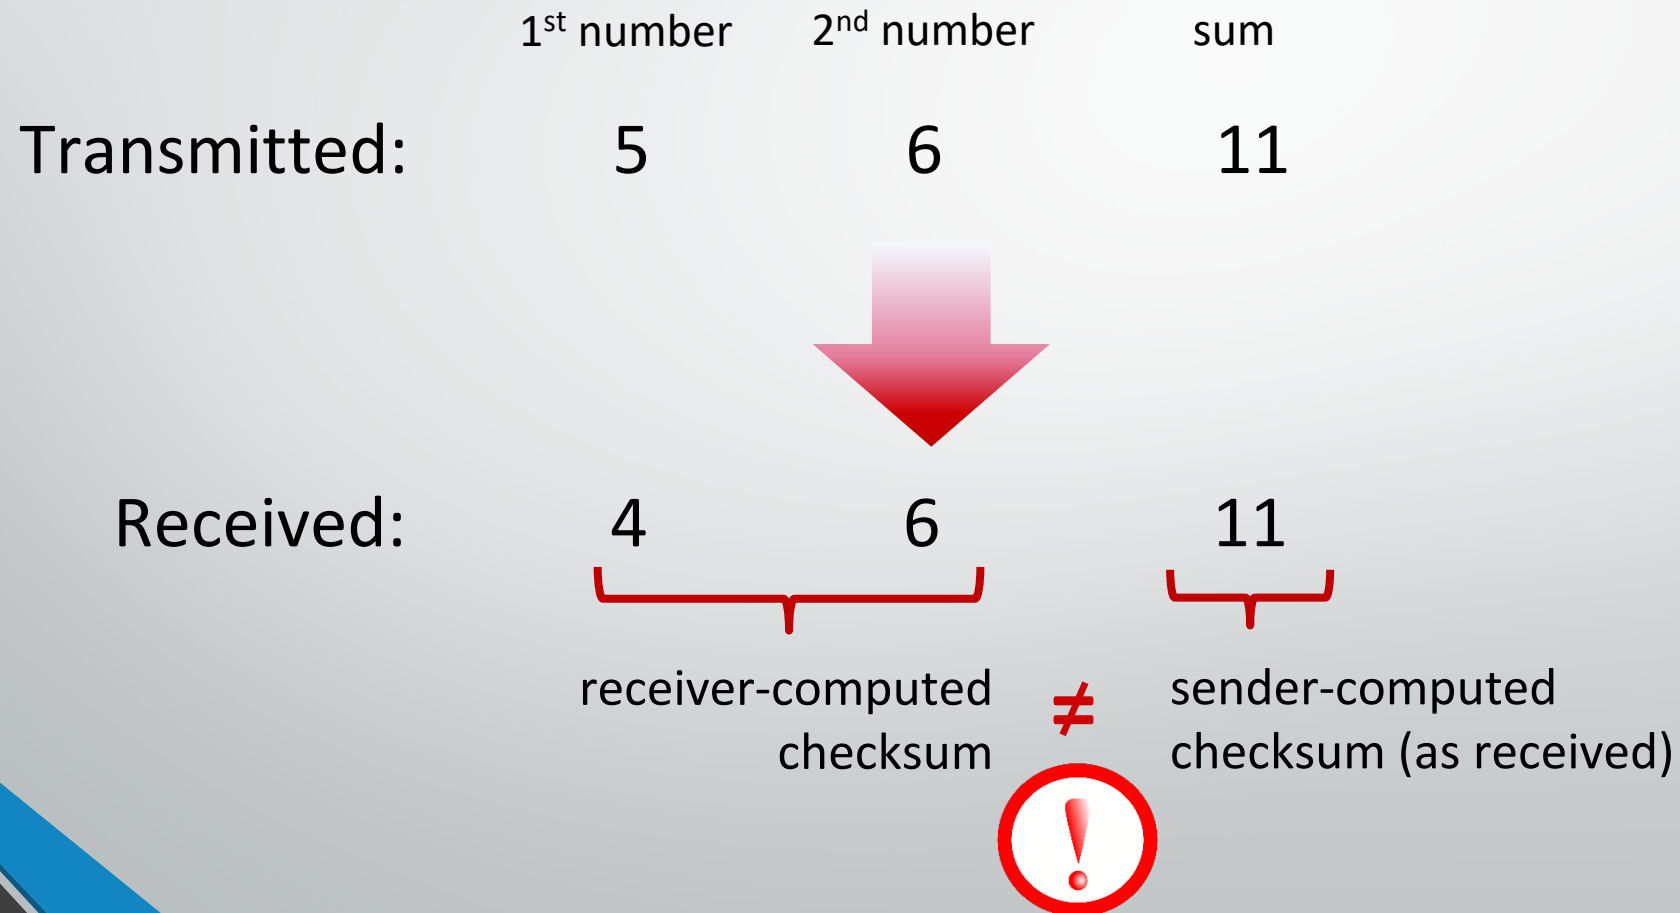


**Figure 14:** User Datagram Header format

- Connectionless
- No reassembly to order.
- No Error checking
- No Flow control

# UDP checksum

*Goal:* detect errors (*i.e.*, flipped bits) in transmitted segment



# UDP checksum

*Goal:* detect errors (*i.e.*, flipped bits) in transmitted segment

## sender:

- treat contents of UDP segment (including UDP header fields and IP addresses) as sequence of 16-bit integers
- **checksum:** addition (one's complement sum) of segment content
- checksum value put into UDP checksum field

## receiver:

- compute checksum of received segment
- check if computed checksum equals checksum field value:
  - Not equal - error detected
  - Equal - no error detected. *But maybe errors nonetheless? More later ....*

# Internet checksum: an example

example: add two 16-bit integers

	1	1	1	0	0	1	1	0	0	1	1	0	0	1	1	0	
	1	1	0	1	0	1	0	1	0	1	0	1	0	1	0	1	
<hr/>																	
wraparound	1	1	0	1	1	1	0	1	1	1	0	1	1	1	0	1	1
<hr/>																	
sum	1	0	1	1	1	0	1	1	1	0	1	1	1	1	0	0	
checksum	0	1	0	0	0	1	0	0	0	1	0	0	0	0	1	1	

*Note:* when adding numbers, a carryout from the **most significant bit** needs to be **added** to the result



# Internet checksum: weak protection!

example: add two 16-bit integers

	1	1	1	0	0	1	1	0	0	1	1	0	0	1	1	0
	1	1	0	1	0	1	0	1	0	1	0	1	0	1	0	1
<hr/>																
wraparound	1	1	0	1	1	1	0	1	1	1	0	1	1	1	0	1
sum	1	0	1	1	1	0	1	1	1	0	1	1	1	1	0	0
checksum	0	1	0	0	0	1	0	0	0	1	0	0	0	0	1	1

Even though numbers have changed (bit flips), *no* change in checksum!



THE END