



Inspiring Excellence

Network Layer

Routing Algorithm
Distance Vector Routing

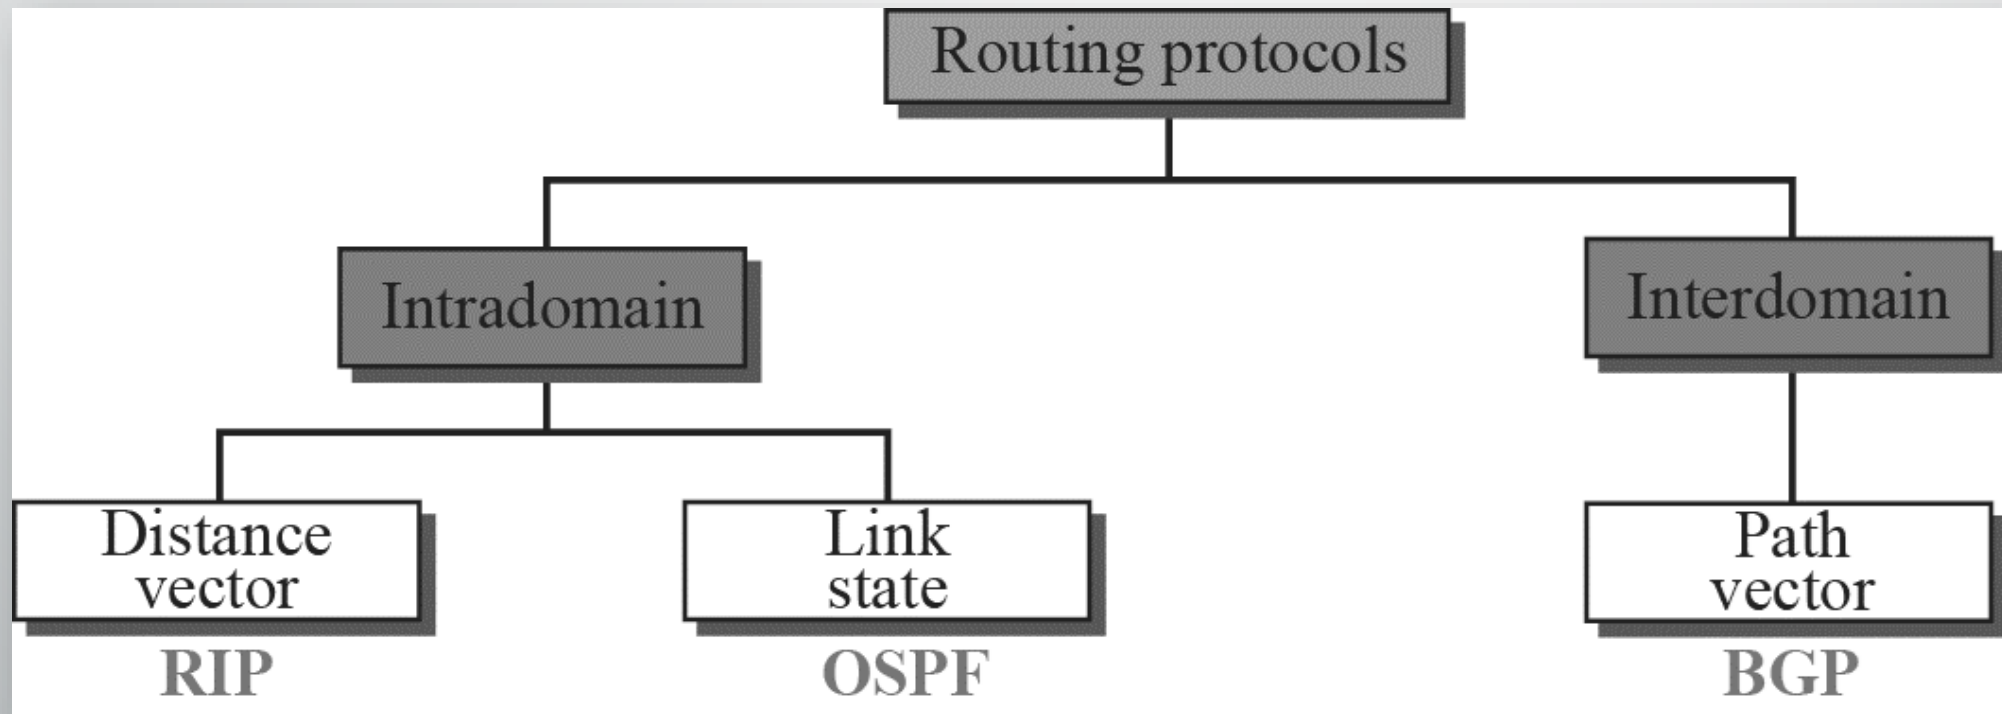
Lecture 13 | Part 1 | CSE421 – Computer Networks

Department of Computer Science and Engineering
School of Data & Science

Objectives

- understand principles behind network layer services:
 - network layer service models
 - forwarding versus routing
 - how a router works
 - routing algorithms
 - **distance vector**
 - **link state**
 - hierarchical routing
 - broadcast, multicast


Popular Routing Protocols



Autonomous Systems

- Internet is divided into **autonomous** systems.
- An autonomous system (AS) is **a group of** networks and routers under the authority of **a single administration**.
- Routing ***inside*** an autonomous system is called **intra-domain routing**. Routing ***between*** autonomous systems is called **inter-domain routing**.

Autonomous Systems



Routing Algorithms

- Given a set of routers and links connecting the routers.
- Routing algorithm finds a “good” path from the source to destination router.
- Good path = Least cost path

Routing Algorithm classification

Static or dynamic?

Static:

- routes **change slowly** over time
- Manually configured

Dynamic:

- routes change more **quickly**
 - in response to link cost changes

Routing Algorithm classification

Global or Decentralized

Global:

- all routers have complete topology and link cost info
- “link state” algorithms


Decentralized:

- router knows physically-connected neighbors, link costs to neighbors
- iterative process of computation, exchange of info with neighbors
- “distance vector” algorithms

Routing on a Graph

- Essentially a graph theory problem
 - Network is a directed graph; routers are vertices
- Find “best” path between every pair of vertices
 - In the simplest case, best path is the shortest path

Graph abstraction




Graph: $G = (N, E)$

$N = \text{set of routers} = \{ u, v, w, x, y, z \}$

$E = \text{set of links} = \{ (u,v), (u,x), (u,w), (v,x), (v,w), (x,w), (x,y), (w,y), (w,z), (y,z) \}$

Graph abstraction: costs



$c(x,x') = \text{cost of link } (x,x')$

- e.g., $c(w,z) = 5$

Question: What's the least-cost path between u and z ?

Routing algorithm: algorithm that finds least-cost path



Distance Vector Algorithm

Distance Vector Algorithm

Distributed:

- each node receives info from one or more of its directly connected neighbors
- Performs a calculations
- Distributes the results back to its neighbors

Iterative

- Process continues until no more info to exchange

Asynchronous:

- All nodes operate independently

Distance Vector Algorithm

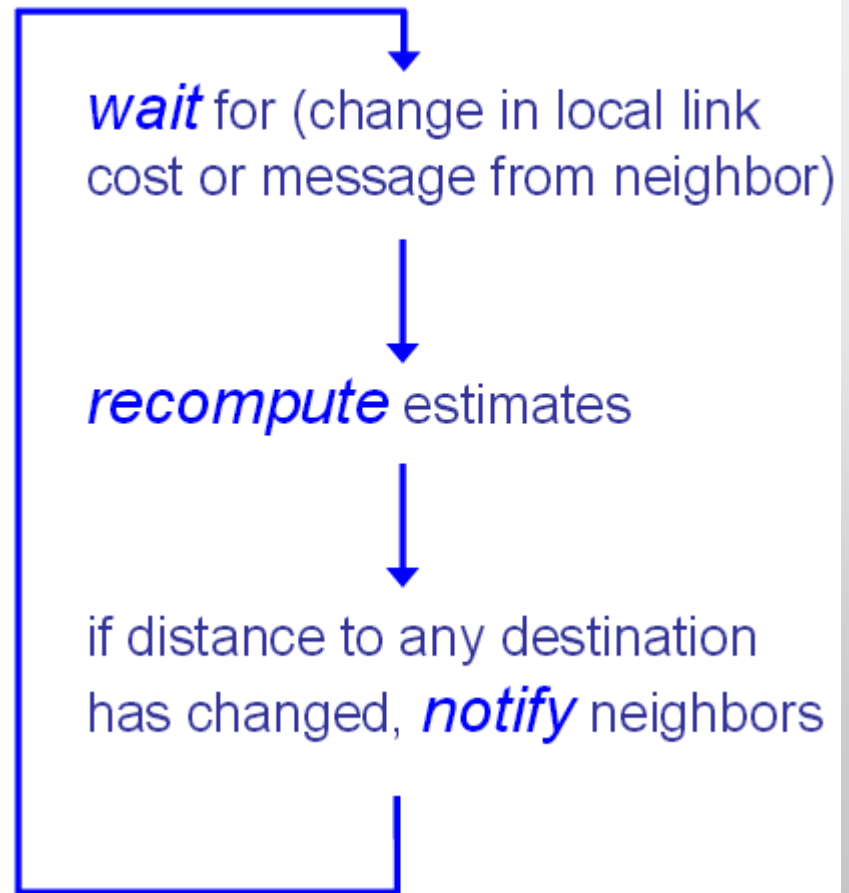
Iterative, asynchronous: each local iteration caused by:

- Local link cost change
- Distance vector update message from neighbor

Distributed:

- Each node notifies neighbors *only* when its DV changes
- Neighbors then notify their neighbors if necessary

Each node:



Distance Vector Algorithm

Bellman-Ford Equation Algorithm

□ computes shortest paths from a single source vertex to all of the other vertices in a weighted digraph.

□ Distributed route computation using only neighbor's info

Distance Vector Algorithm

Objective:

$d_x(y) :=$ cost of least-cost path from x to y

Then

$$d_x(y) = \min_v \{c(x,v) + d_v(y)\}$$


where min is taken over all neighbors v of x

Distance Vector Algorithm

With the Distance Vector Routing algorithm, the node x contains the following routing information:

- For each neighbor v , the cost $c(x,v)$ is the path cost from x to directly attached neighbor, v .
- The distance vector x , i.e., $D_x = [D_x(y) : y \text{ in } N]$, containing its cost to all destinations, y , in N .
- The distance vector of each of its neighbors, i.e., $D_v = [D_v(y) : y \text{ in } N]$ for each neighbor v of x .

Bellman-Ford example from u to z



U has 3 neighbors v, x and w

Clearly, $d_v(z) = 5$, $d_x(z) = 3$, $d_w(z) = 3$

B-F equation says:

$$\begin{aligned} d_u(z) &= \min \{ c(u,v) + d_v(z), \\ &\quad c(u,x) + d_x(z), \\ &\quad c(u,w) + d_w(z) \} \\ &= \min \{ 2 + 5, \\ &\quad 1 + 3, \\ &\quad 5 + 3 \} = 4 \end{aligned}$$

- ❑ Node that achieves minimum is the next hop in shortest path to a destination,
- ❑ To go to z from u, x is the next hop in the forwarding table

Distance vector algorithm

Basic idea:

- Each node periodically sends its own distance vector estimate to neighbors
- When a node x receives new DV estimate from neighbor;
- It updates its own DV using B-F equation:

$$D_x(y) = \min\{c(x,y) + D_y(y), c(x,z) + D_z(y)\}$$

$$= \min\{2+0, 7+1\} = 2$$

$$D_x(z) = \min\{c(x,y) + D_y(z), c(x,z) + D_z(z)\}$$

$$= \min\{2+1, 7+0\} = 3$$

node x table

		cost to		
		x	y	z
from	x	0	2	7
	y	∞	∞	∞
	z	∞	∞	∞


node y table

		cost to		
		x	y	z
from	x	∞	∞	∞
	y	2	0	1
	z	∞	∞	∞

node z table

		cost to		
		x	y	z
from	x	∞	∞	∞
	y	∞	∞	∞
	z	7	1	0

		cost to		
		x	y	z
from	x	0	2	3
	y	2	0	1
	z	7	1	0



► time

$$Dy(x) = \min\{c(y,x) + Dx(x), c(y,z) + D_z(x)\}$$

$$= \min\{2+0, 1+7\} = 2$$

$$Dy(z) = \min\{c(y,x) + Dx(z), c(y,z) + D_z(z)\}$$

$$= \min\{2+7, 1+0\} = 1$$

node x table

		cost to		
		x	y	z
from	x	0	2	7
	y	∞	∞	∞
	z	∞	∞	∞

node y table


		cost to		
		x	y	z
from	x	∞	∞	∞
	y	2	0	1
	z	∞	∞	∞

node z table

		cost to		
		x	y	z
from	x	∞	∞	∞
	y	∞	∞	∞
	z	7	1	0

		cost to		
		x	y	z
from	x	0	2	3
	y	2	0	1
	z	7	1	0

		cost to		
		x	y	z
from	x	0	2	7
	y	2	0	1
	z	7	1	0



time

$$Dz(x) = \min\{c(z,x) + Dx(x), c(z,y) + Dy(x)\}$$

$$= \min\{7+0, 1+2\} = 3$$

$$Dz(y) = \min\{c(z,x) + Dx(y), c(z,y) + Dy(y)\}$$

$$= \min\{7+2, 1+0\} = 1$$

node x table

		cost to		
		x	y	z
from	x	0	2	7
	y	∞	∞	∞
	z	∞	∞	∞

node y table

		cost to		
		x	y	z
from	x	∞	∞	∞
	y	2	0	1
	z	∞	∞	∞


node z table

		cost to		
		x	y	z
from	x	∞	∞	∞
	y	∞	∞	∞
	z	7	1	0

		cost to		
		x	y	z
from	x	0	2	3
	y	2	0	1
	z	7	1	0

		cost to		
		x	y	z
from	x	0	2	7
	y	2	0	1
	z	7	1	0

		cost to		
		x	y	z
from	x	0	2	7
	y	2	0	1
	z	3	1	0



► time

node x table

		cost to		
		x	y	z
from	x	0	2	7
	y	∞	∞	∞
	z	∞	∞	∞

node y table

		cost to		
		x	y	z
from	x	∞	∞	∞
	y	2	0	1
	z	∞	∞	∞

node z table

		cost to		
		x	y	z
from	x	∞	∞	∞
	y	∞	∞	∞
	z	7	1	0

		cost to		
		x	y	z
from	x	0	2	3
	y	2	0	1
	z	7	1	0


		cost to		
		x	y	z
from	x	0	2	7
	y	2	0	1
	z	7	1	0

		cost to		
		x	y	z
from	x	0	2	7
	y	2	0	1
	z	3	1	0

		cost to		
		x	y	z
from	x	0	2	3
	y	2	0	1
	z	3	1	0


		cost to		
		x	y	z
from	x	0	2	3
	y	2	0	1
	z	3	1	0

		cost to		
		x	y	z
from	x	0	2	3
	y	2	0	1
	z	3	1	0



time

Example



- How many iterations are needed to make the final routing tables of each router?
- What will be the routing table of B and D after the 2nd iteration?



Inspiring Excellence

Network Layer

Routing Algorithm
Distance Vector Routing

Lecture 13 | Part 2 | CSE421 – Computer Networks

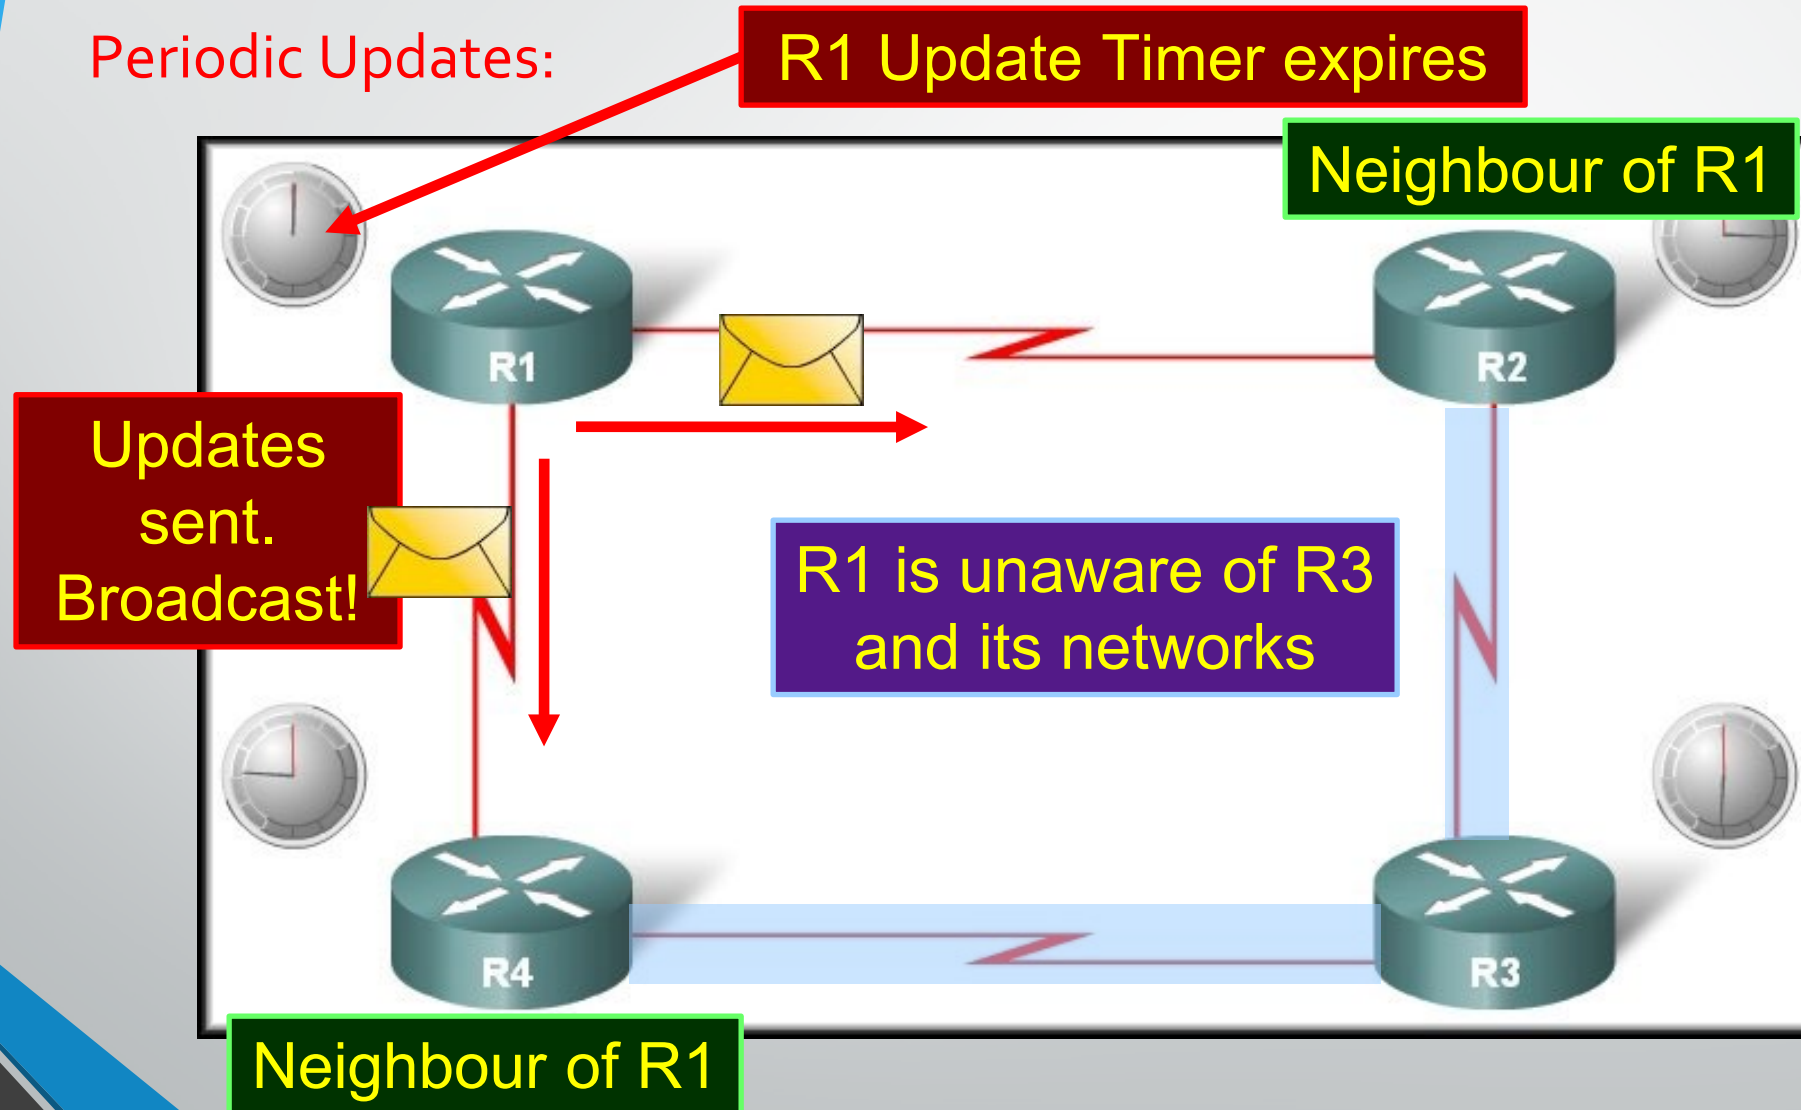
Department of Computer Science and Engineering
School of Data & Science



Operation of Distance Vector

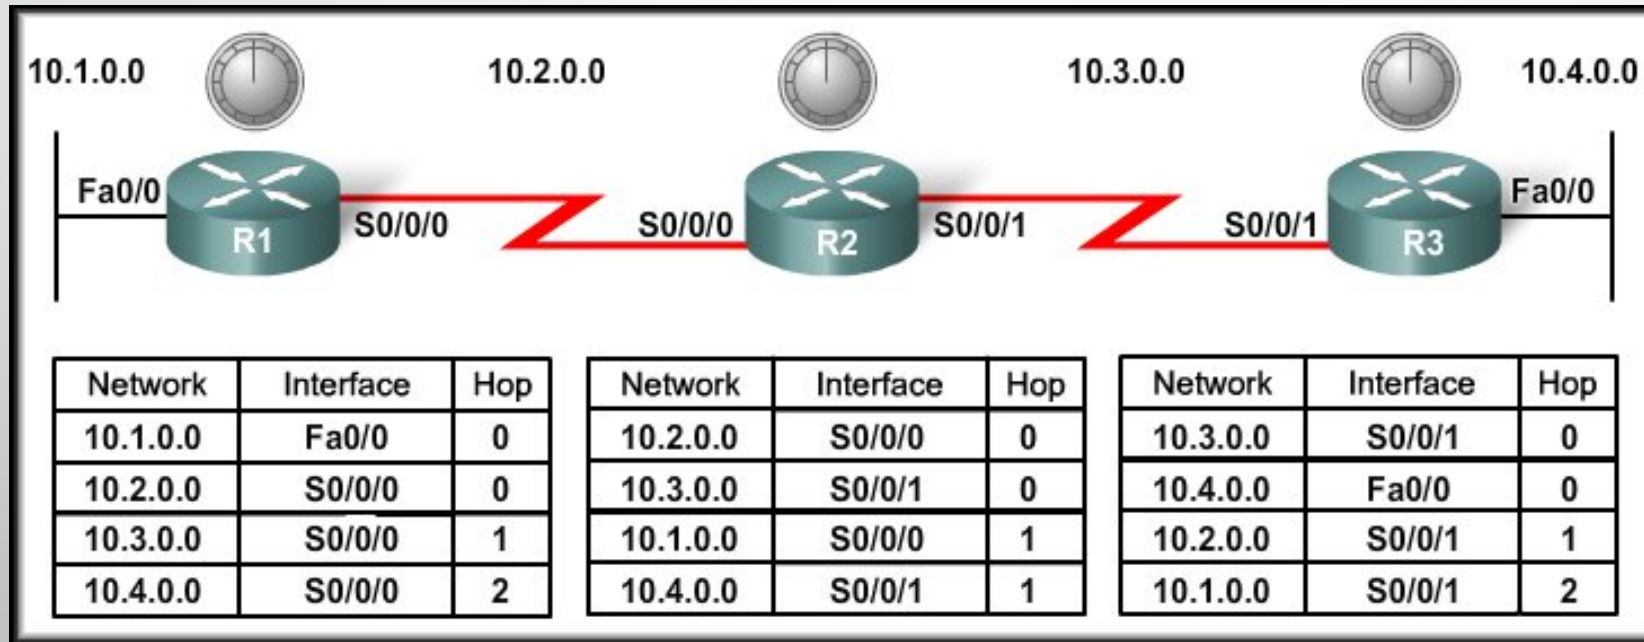
- Periodic Updates:
 - Periodically broadcast the entire routing table to each of its neighbors (RIP – every 30 seconds).
 - Inefficient
- Router is only aware of the:
 - Network addresses of its own interfaces.
 - Network addresses the neighbors running the same routing protocol.


Operation of Distance Vector



Distance Vector Routing Protocols

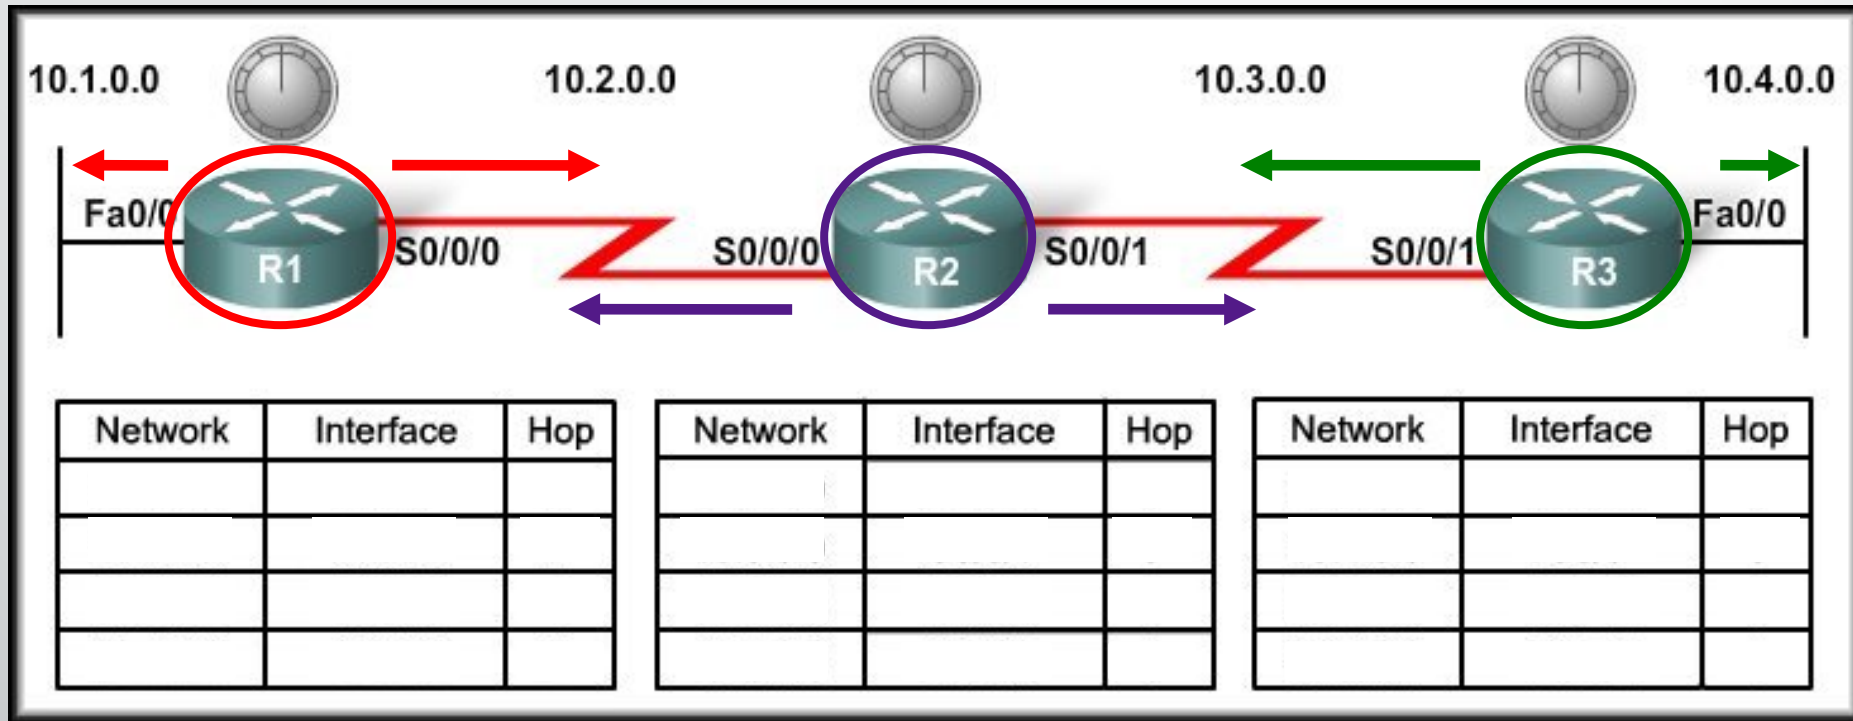
Network Discovery





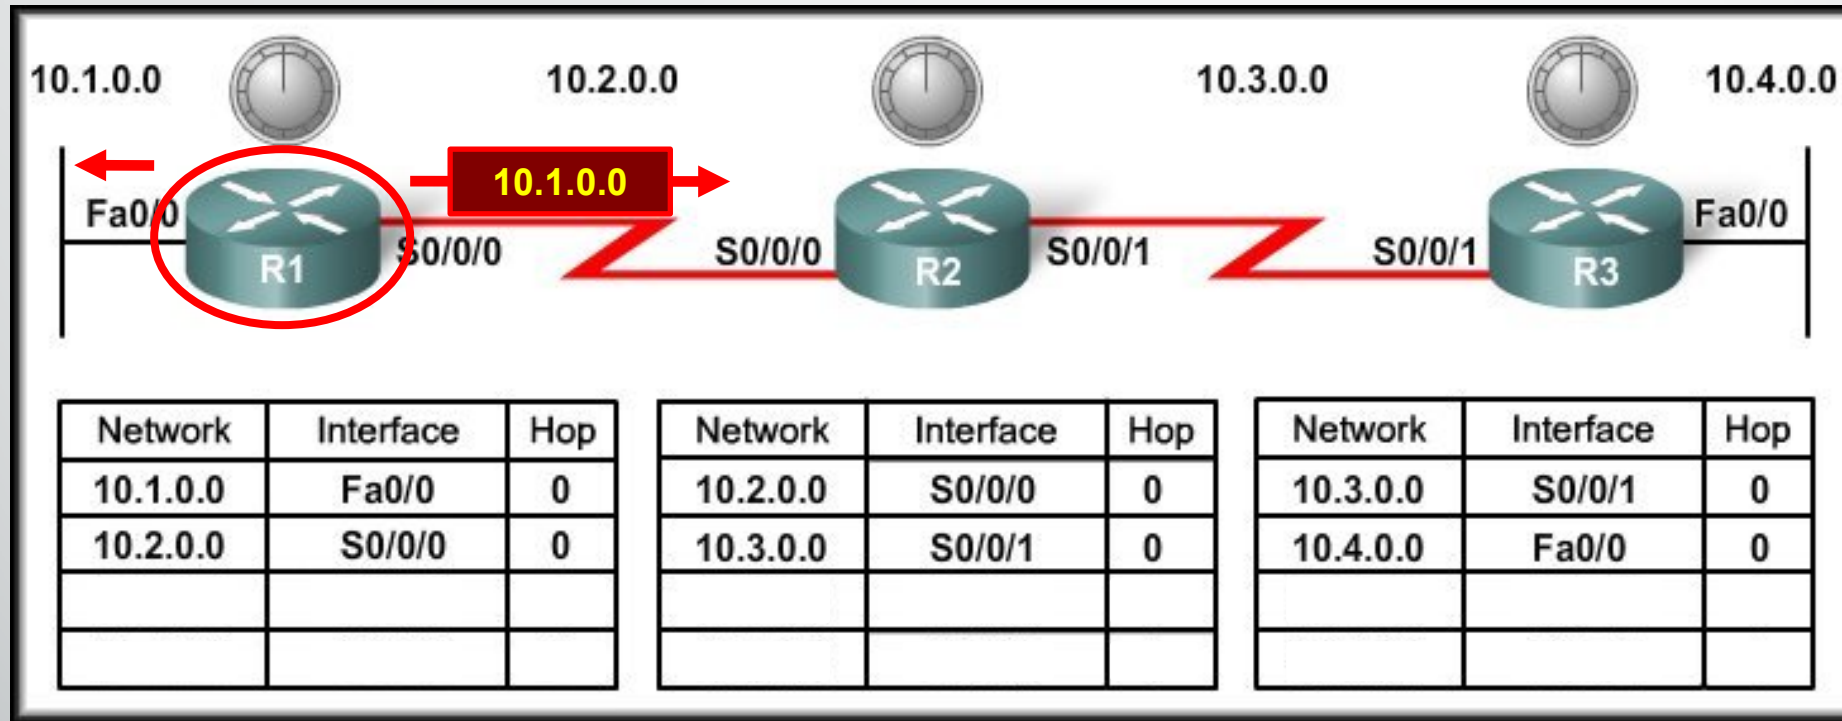
- Network Discovery:
 - Is part of the process of the routing protocol algorithm that enables routers to **learn about remote networks for the first time.**

Cold Start



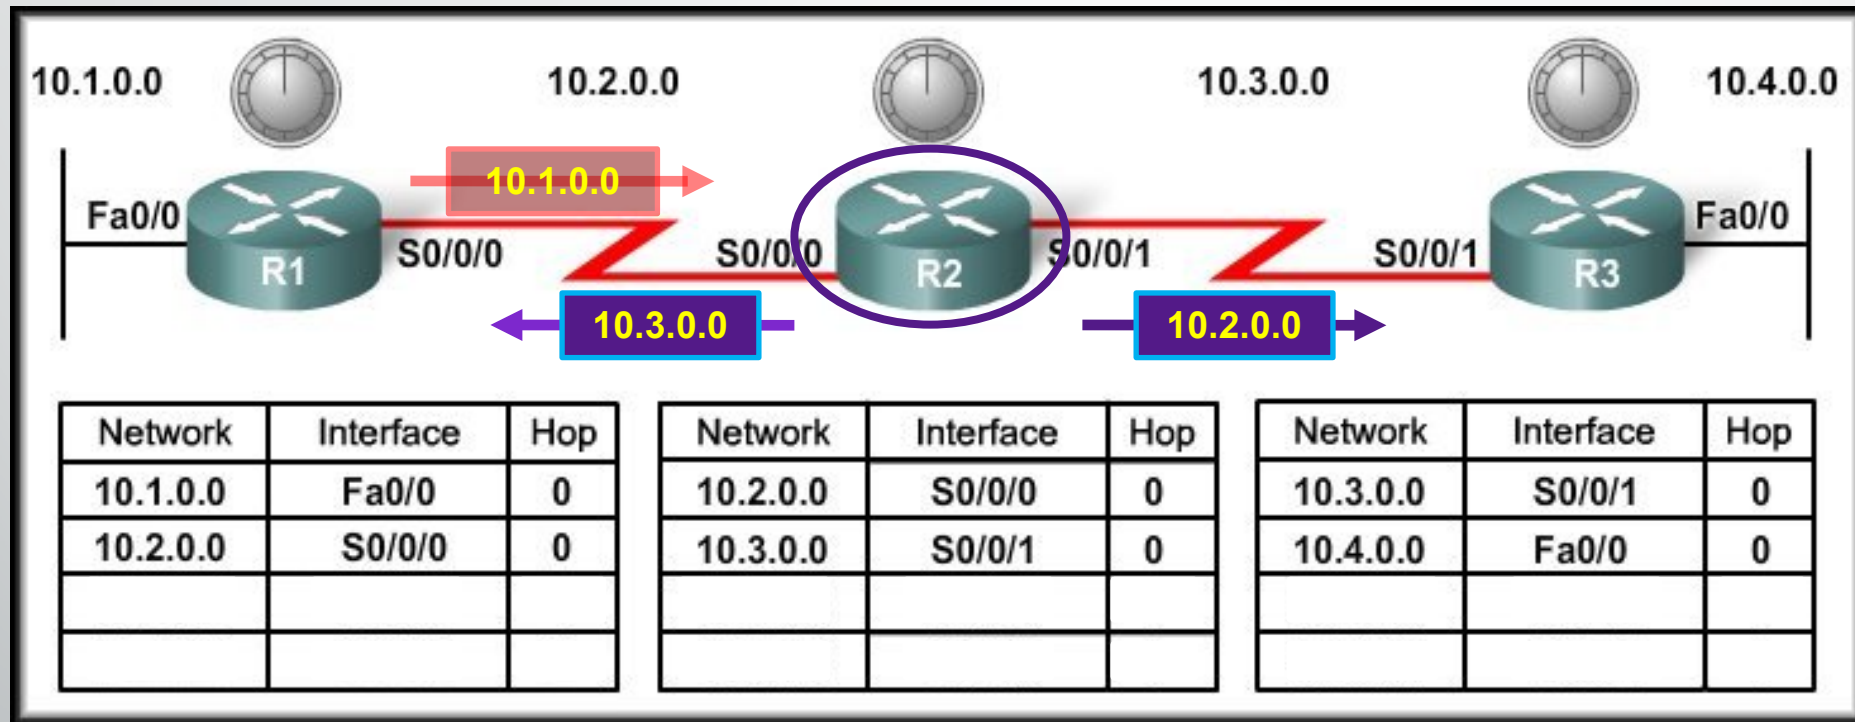
- When a router powers up:
 - Knows nothing about the network topology.
 - Knows only the information saved in NVRAM.
 - Sends updates about its known networks out all ports.

Initial Exchange of Routing Information

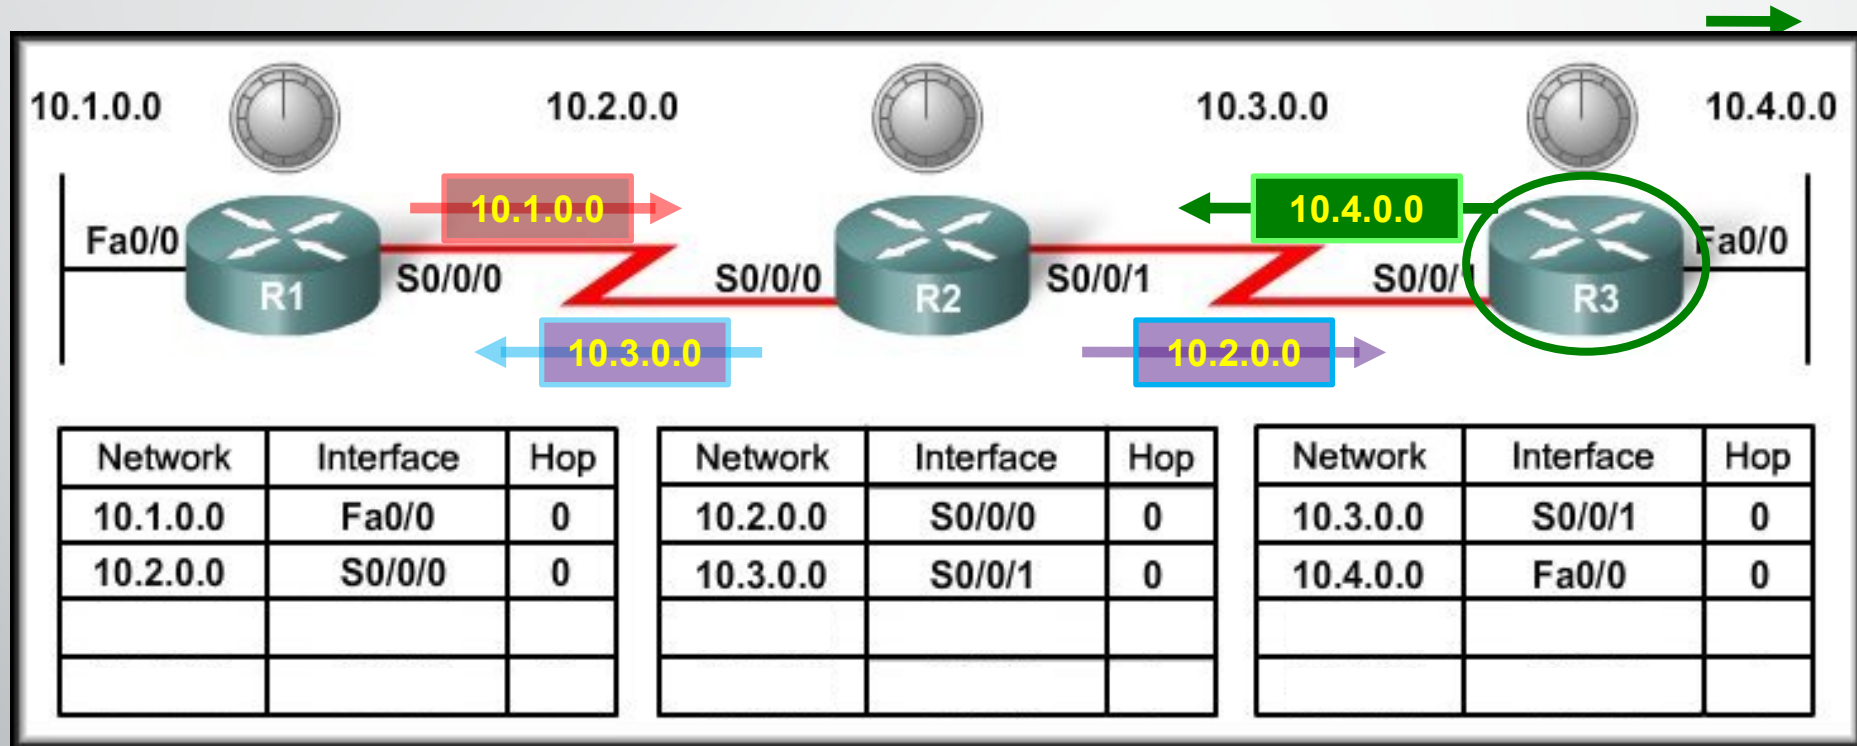


- Sends an update about network **10.1.0.0** out the **Serial o/o/o** interface with a metric of 1.
- Sends an update about network **10.2.0.0** out the **Fao/o** interface with a metric of 1.

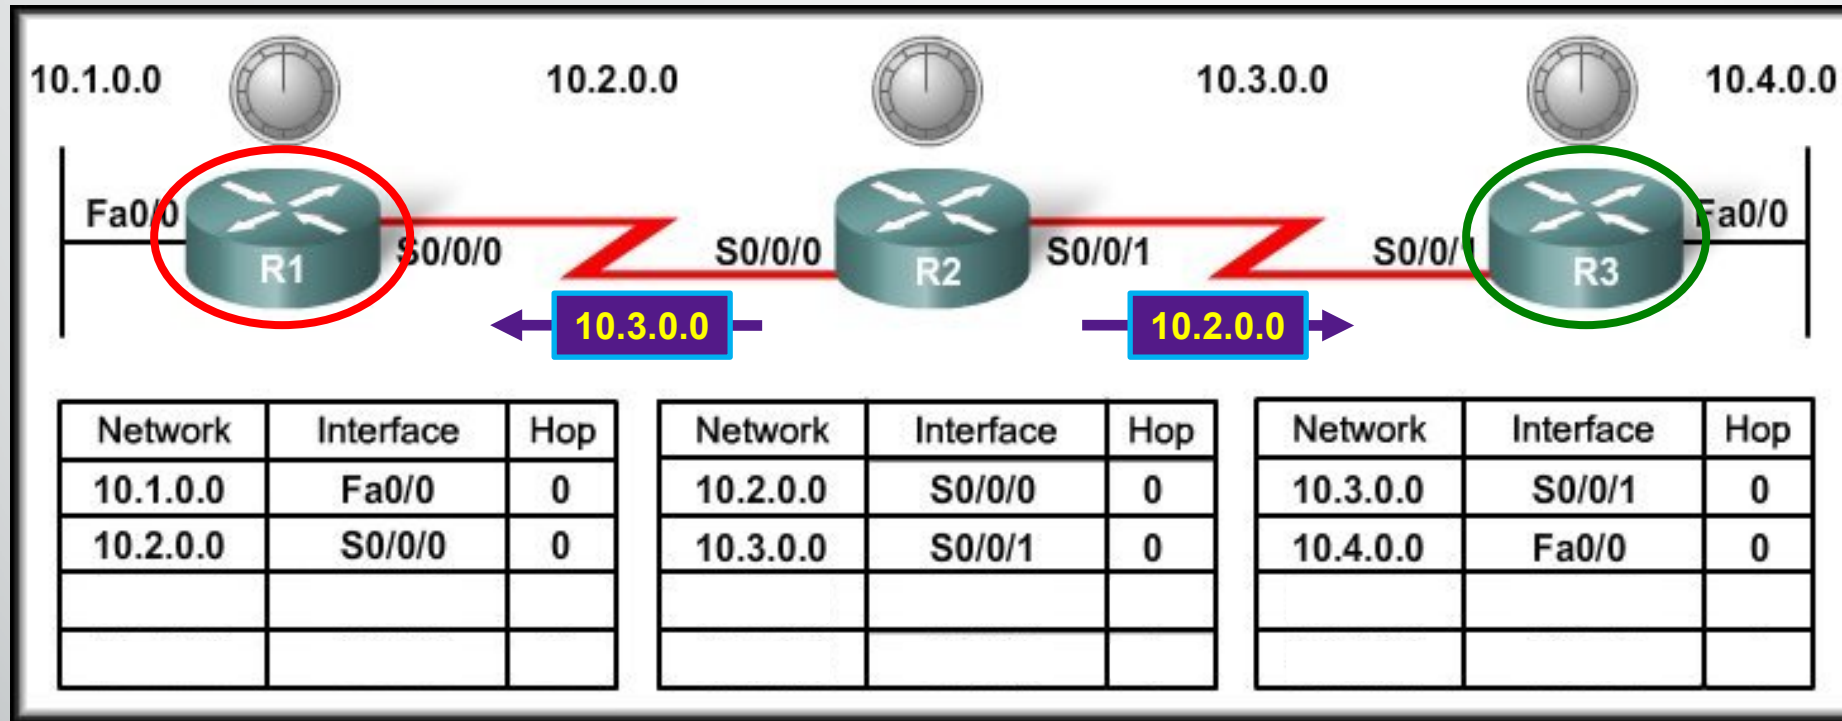
Initial Exchange of Routing Information



Initial Exchange of Routing Information

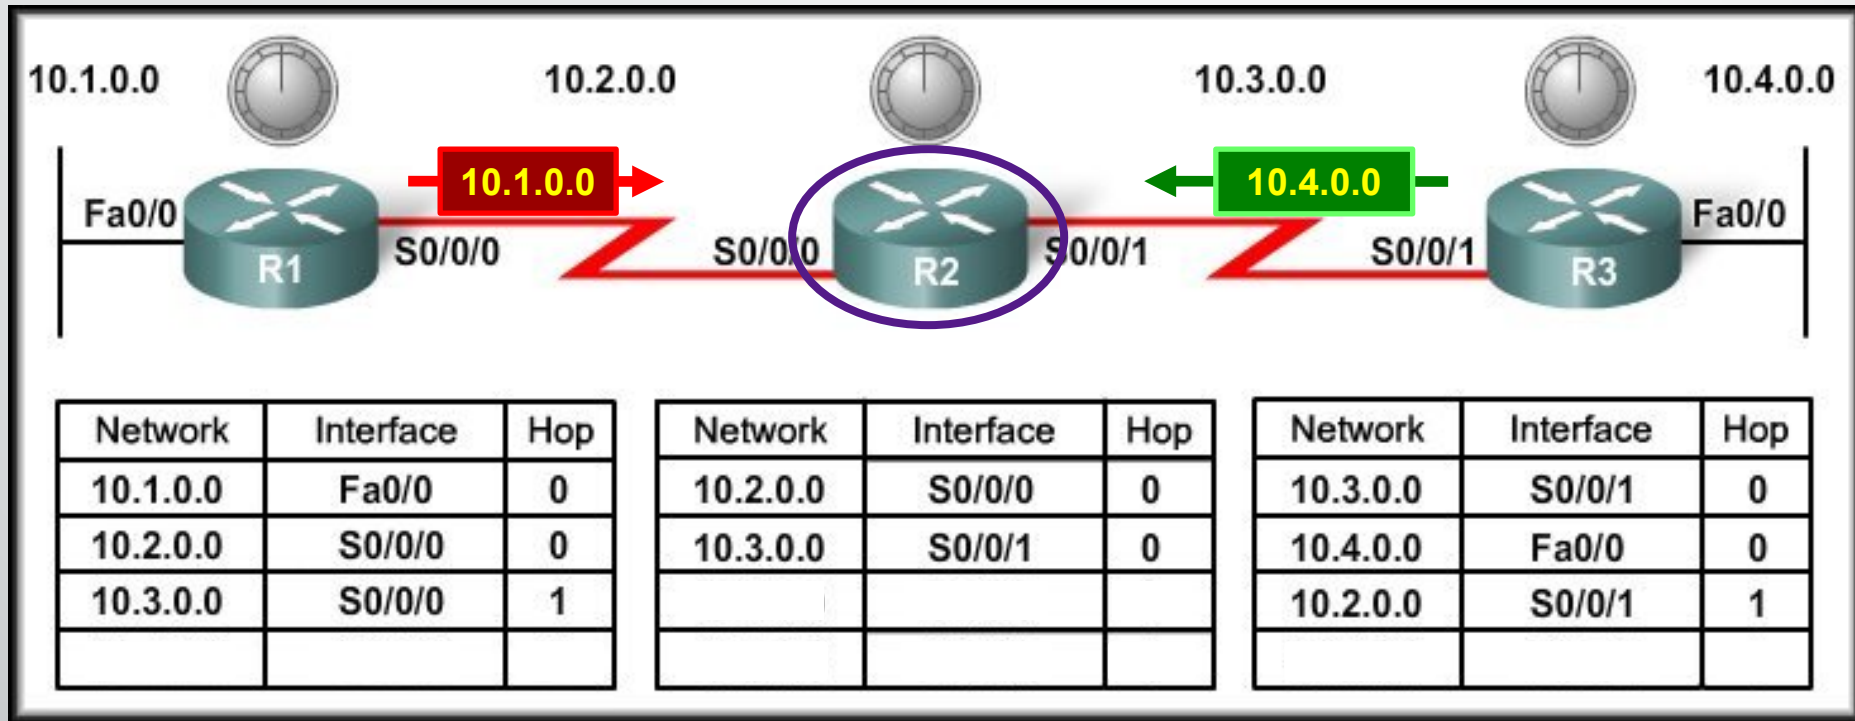


Initial Exchange of Routing Information



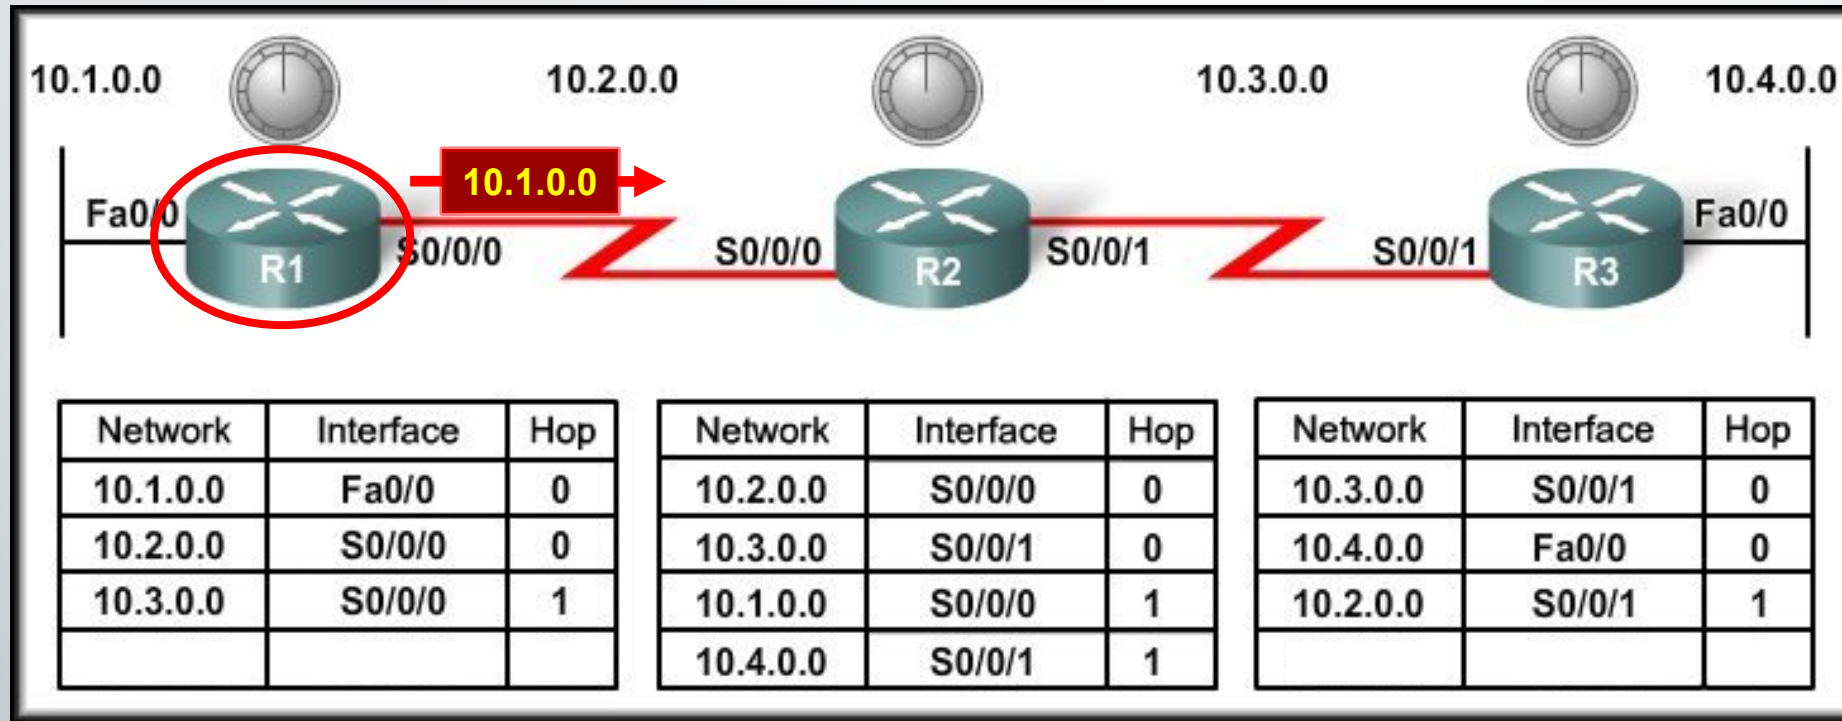
- **R1 Receives** the update from R2 about network 10.3.0.0 and adds it to its routing table.
- **R3 Receives** the update from R2 about network 10.2.0.0 and adds it to its routing table.

Initial Exchange of Routing Information



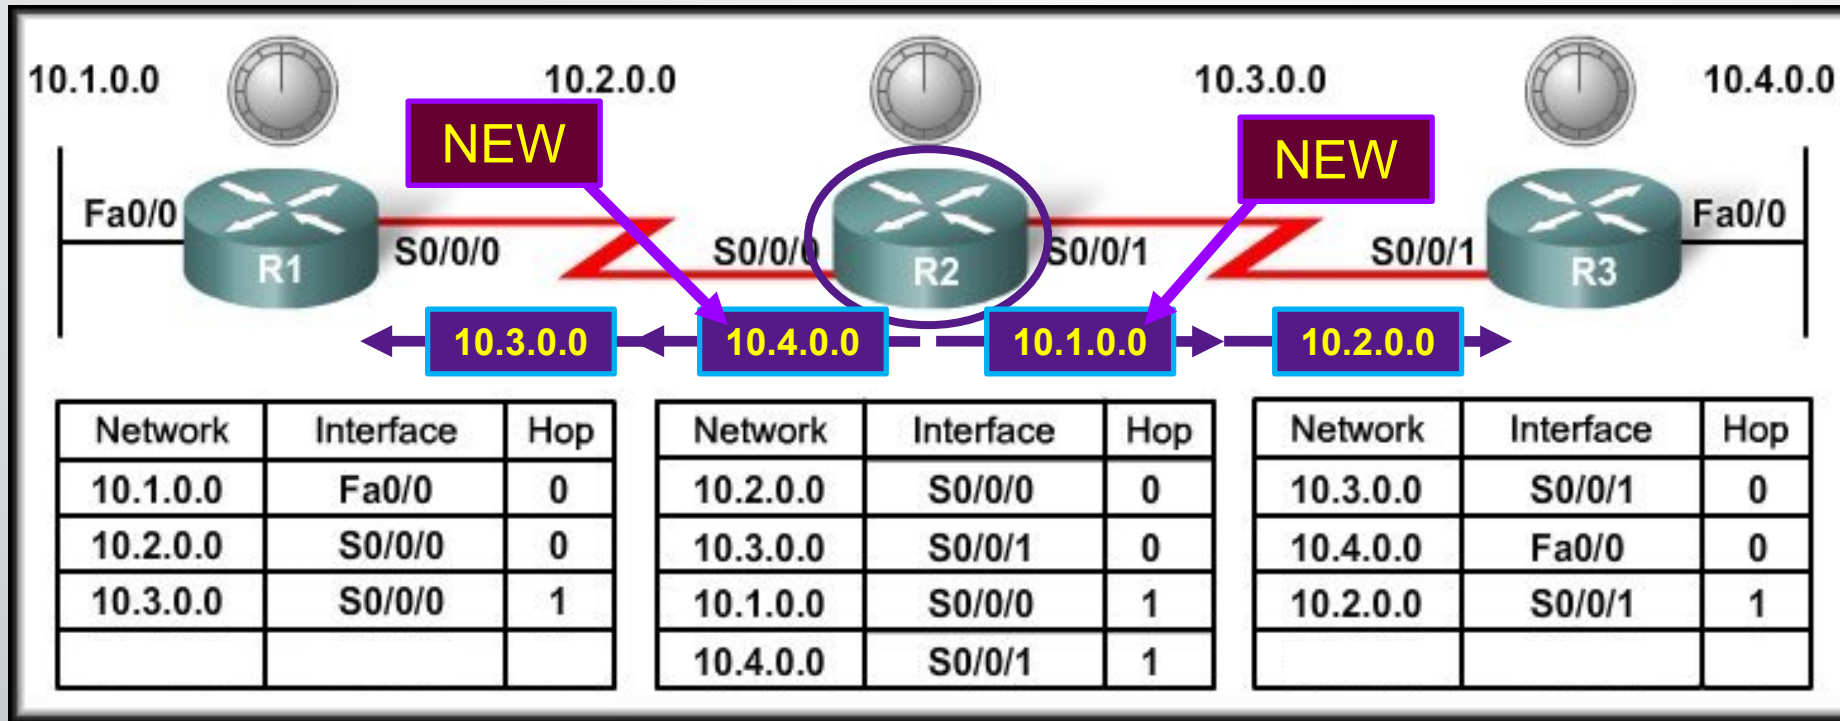
- R2 Receives the update from R1 about network 10.1.0.0 and adds it to its routing table.
- R2 Receives the update from R3 about network 10.4.0.0 and adds it to its routing table.

Next Exchange of Routing Information



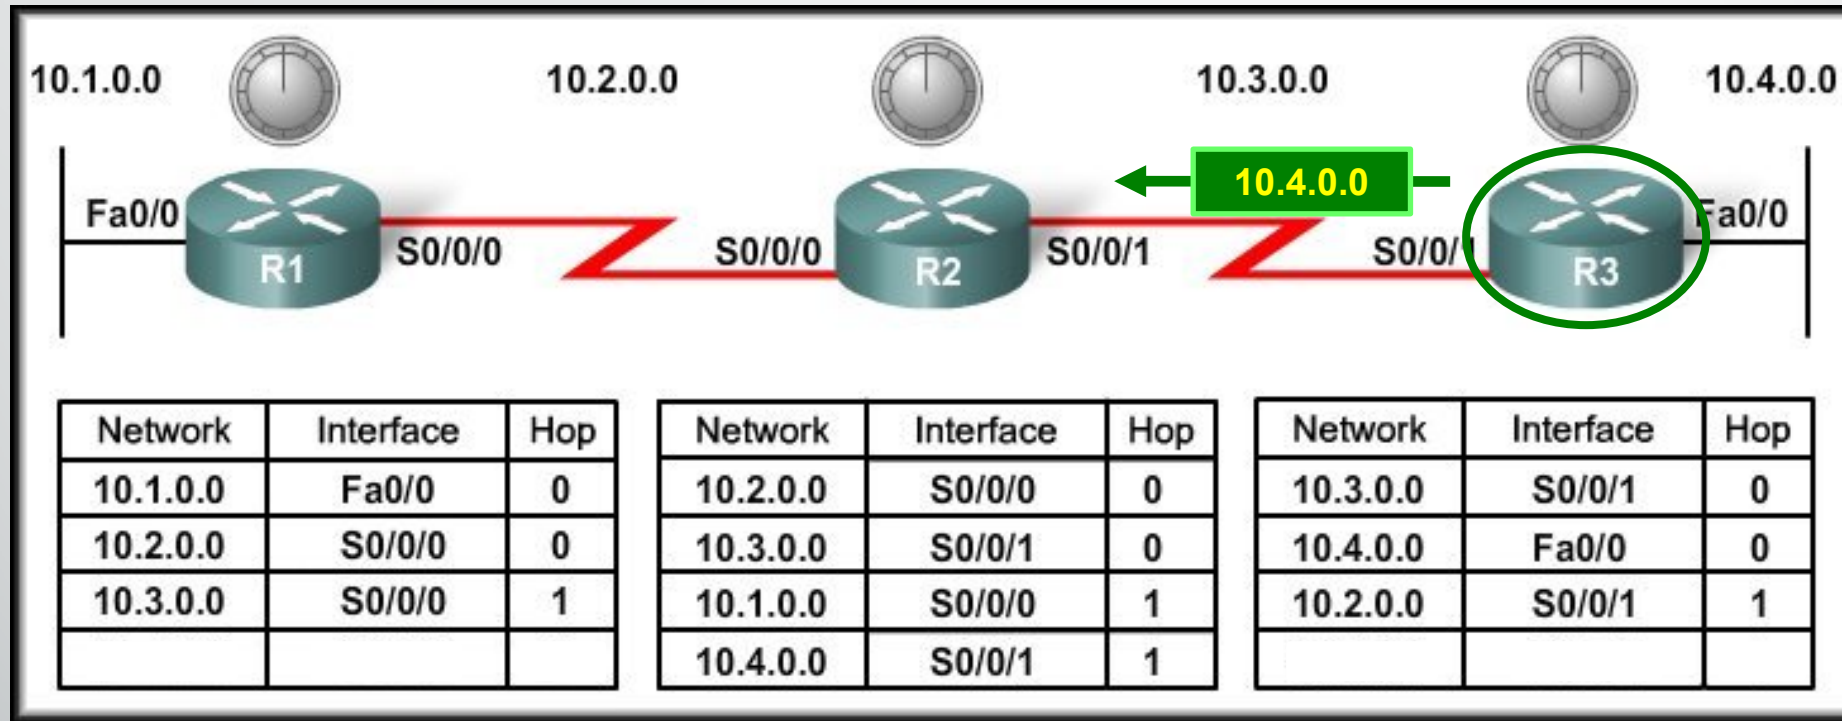
- Sends an update about network **10.1.0.0** out the S0/0/0 interface with a metric of 1 - **AGAIN!**
- **When R2 receives the update**, there is **no change** in information so the update is ignored.

Next Exchange of Routing Information



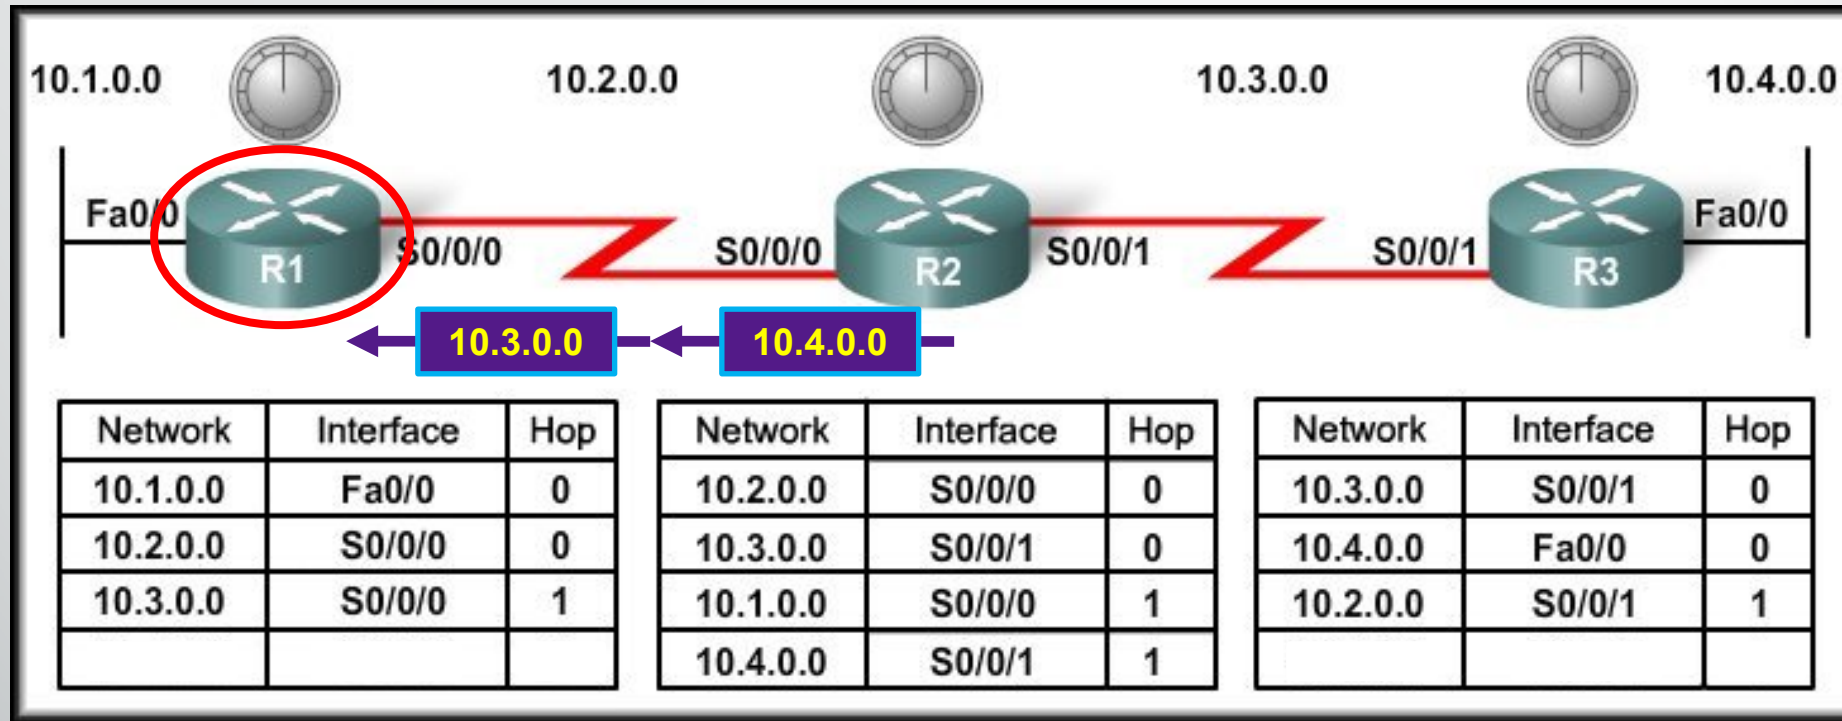
- Sends an update about networks **10.3.0.0** with a metric of 1 and **10.4.0.0** with a metric of 2 out the **Serial o/o/o** interface.
- Similarly sends updates about networks **10.1.0.0** with a metric of 2 and **10.2.0.0** with a metric of 1 out the **Serial o/o/1** interface.

Next Exchange of Routing Information



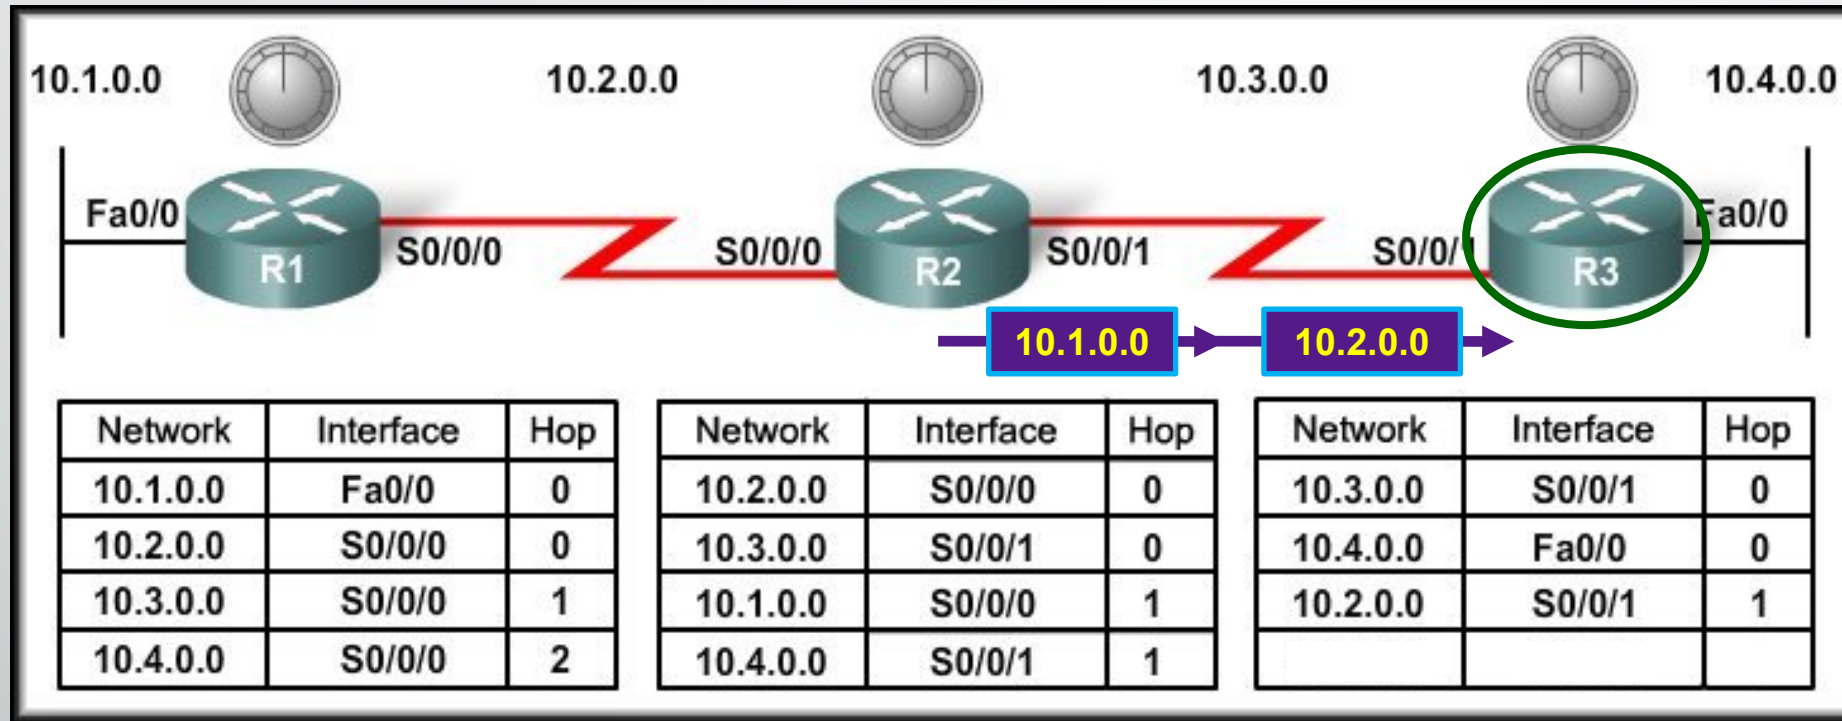
- r Sends an update about network 10.4.0.0 out the So/o/o interface with a metric of 1 - **AGAIN!**
- r When R2 receives the update, there is no change in information so the update is ignored.

Next Exchange of Routing Information



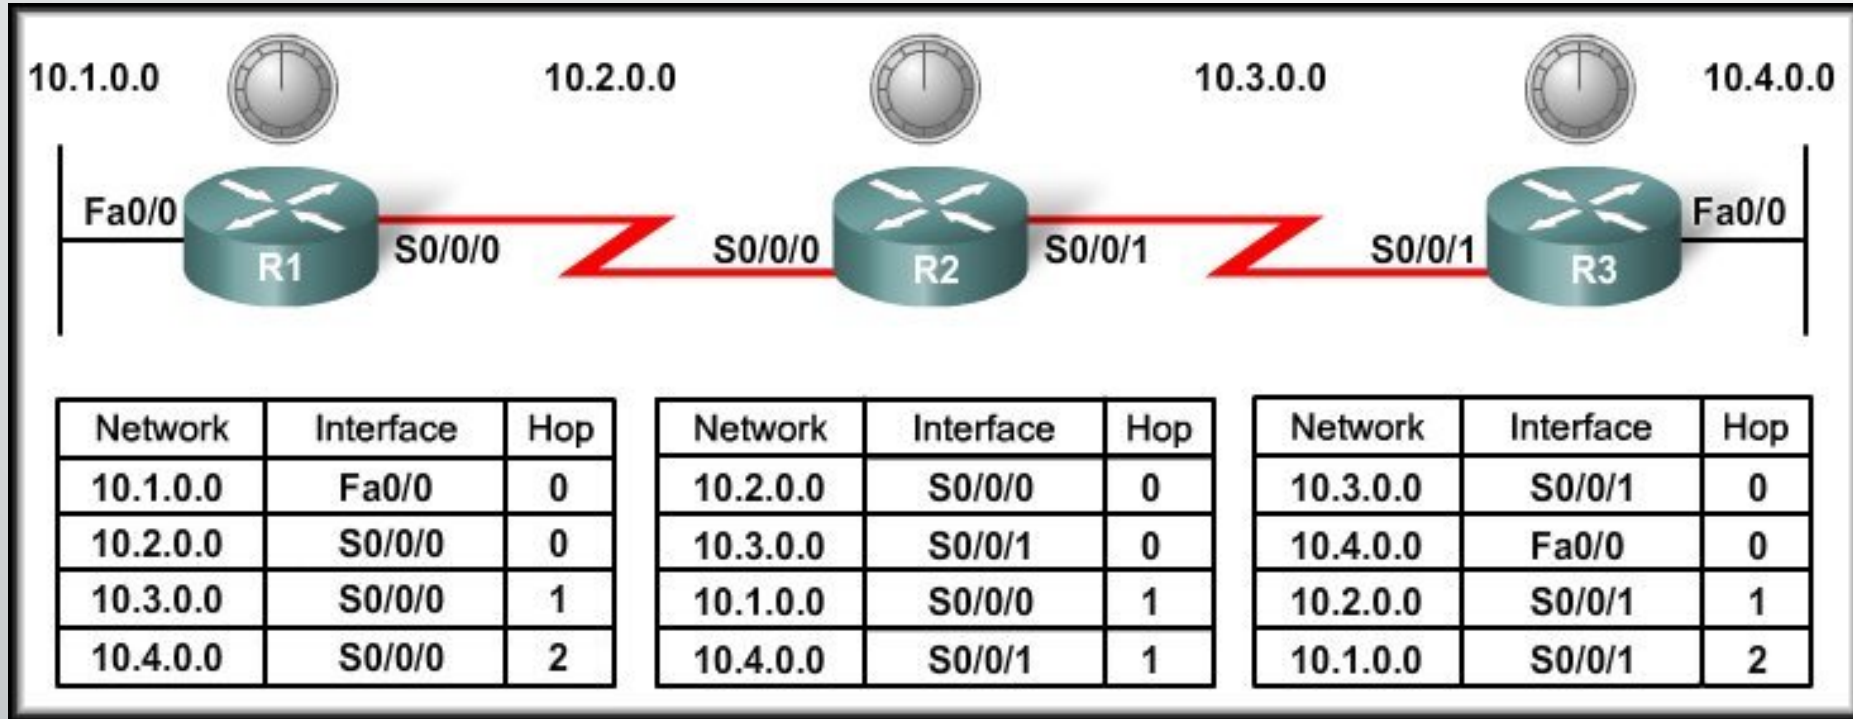
- R1 receives an update from R2 about network 10.3.0.0 and there is no change – update ignored.
- R1 receives an update from R2 about network 10.4.0.0 (new) and adds it to its routing table.

Next Exchange of Routing Information



- R3 receives an update from R2 about network 10.2.0.0 and there is no change – update ignored.
- R3 receives an update from R2 about network 10.1.0.0 (new) and adds it to its routing table.


Next Exchange of Routing Information



- The network has **CONVERGED!**
 - All routers now know about all of the networks attached to all of their neighbouring routers.

Convergence

- The amount of time it takes for a network to converge is directly proportional to the size of that network.
- Routing protocols are compared based on how fast they can propagate this information - their speed to convergence.



- A network is not completely operable until it has converged.
 - Network administrators prefer routing protocols with shorter convergence times.



THE END